

Homework 3: FPP3 Toolbox Exercises

Randy Howk

February 22, 2026

Contents

1 Exercise	2
2 Exercise	6
2.1 a. Time plot of the series	6
2.2 b. Drift forecasts	7
2.3 c. Show drift forecast equals extension of first-to-last line	7
2.4 d. Other benchmark forecasts and best model	10
3 Exercise	11
4 Exercise	13
5 Exercise (For Fun)	16
6 Exercise	17
7 Exercise	17
8 Exercise	21
8.1 8(d) Residual diagnostics for preferred method	22
9 Exercise	23
10 Exercise	24
11 Exercise	26
12 Exercise	26

1 Exercise

```
aus_pop <- global_economy |>
  filter(Country == "Australia") |>
  select(Year, Population)

bricks <- aus_production |>
  select(Quarter, Bricks) |>
  tidyr::drop_na()

lambs_nsw <- aus_livestock |>
  filter(State == "New South Wales", Animal == "Lambs") |>
  select(Month, Count)

wealth_aus <- hh_budget |>
  filter(Country == "Australia") |>
  select(Year, Wealth)

takeaway_aus <- aus_retail |>
  filter(Industry == "Takeaway food services") |>
  summarise(Turnover = sum(Turnover))
```

```
interval(aus_pop)
```

```
## <interval[1]>
## [1] 1Y
```

```
interval(bricks)
```

```
## <interval[1]>
## [1] 1Q
```

```
interval(lambs_nsw)
```

```
## <interval[1]>
## [1] 1M
```

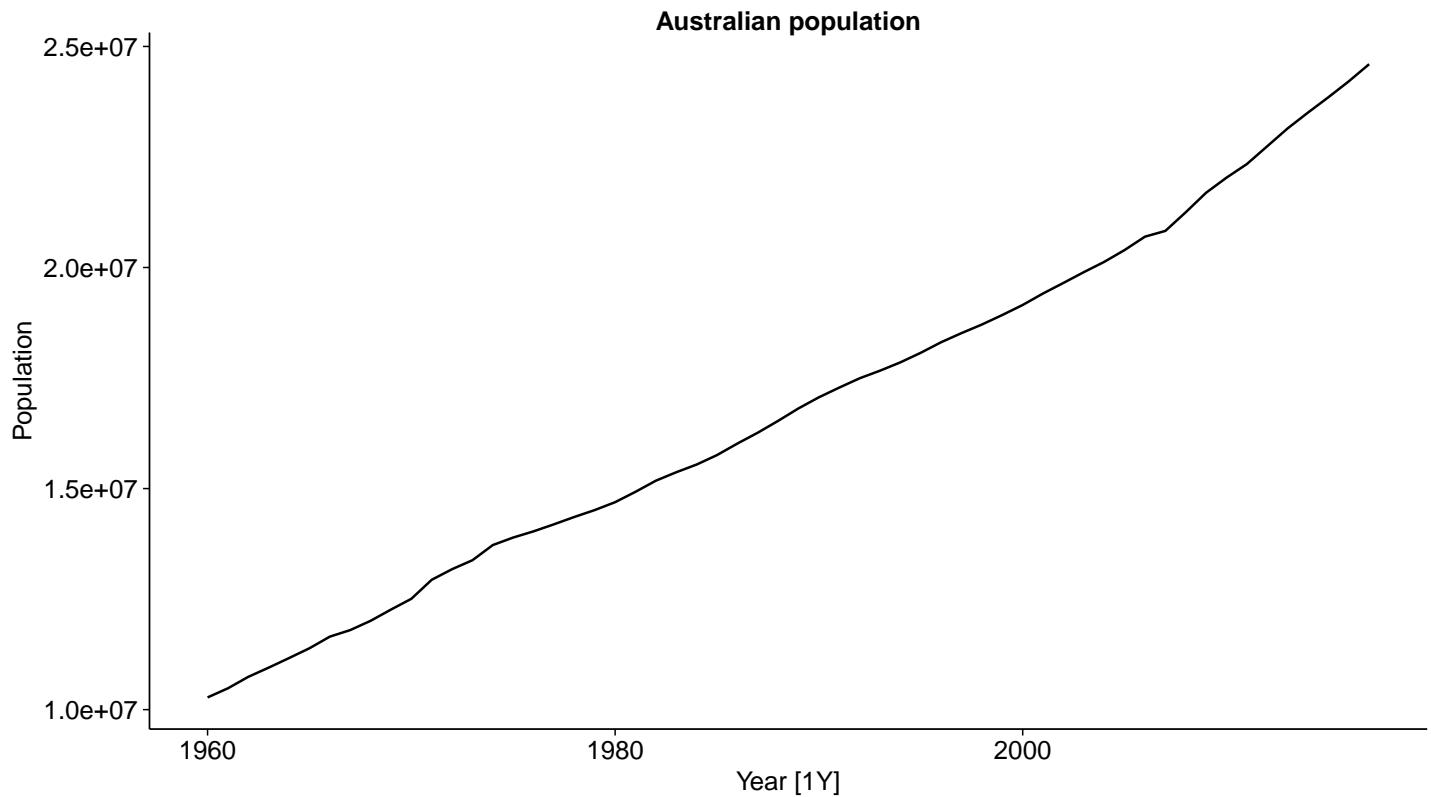
```
interval(wealth_aus)
```

```
## <interval[1]>
## [1] 1Y
```

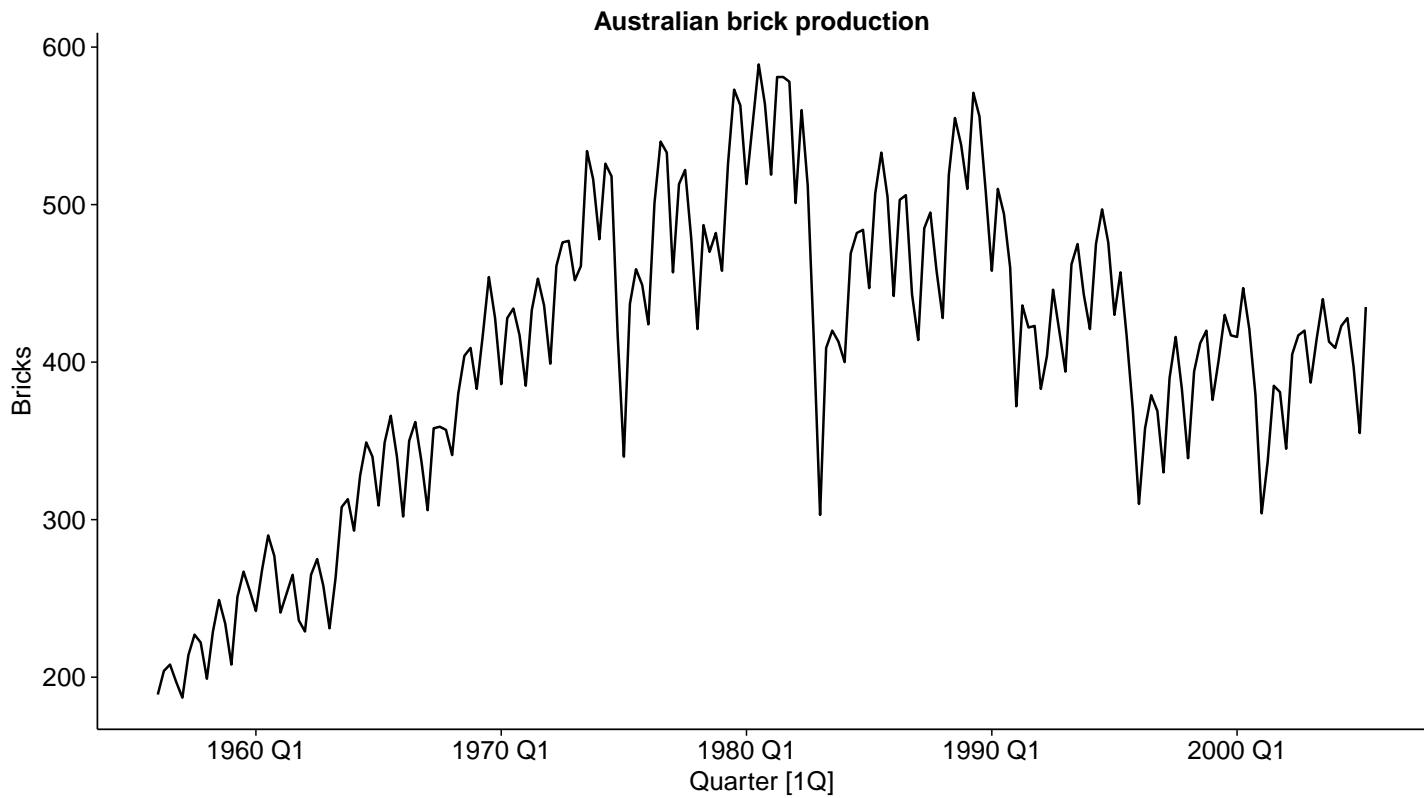
```
interval(takeaway_aus)
```

```
## <interval[1]>
## [1] 1M
```

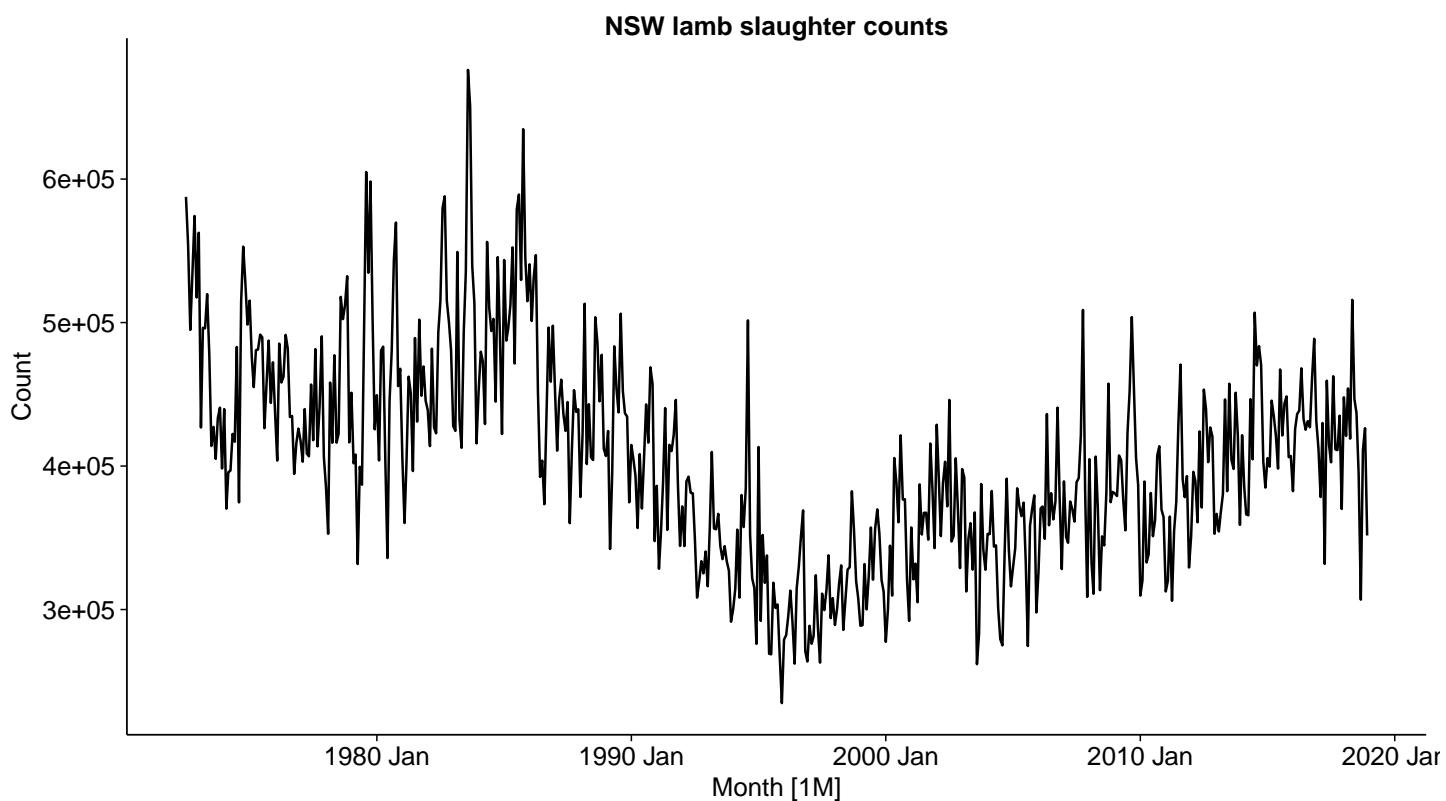
```
aus_pop |> autoplot(Population) + labs(title = "Australian population")
```



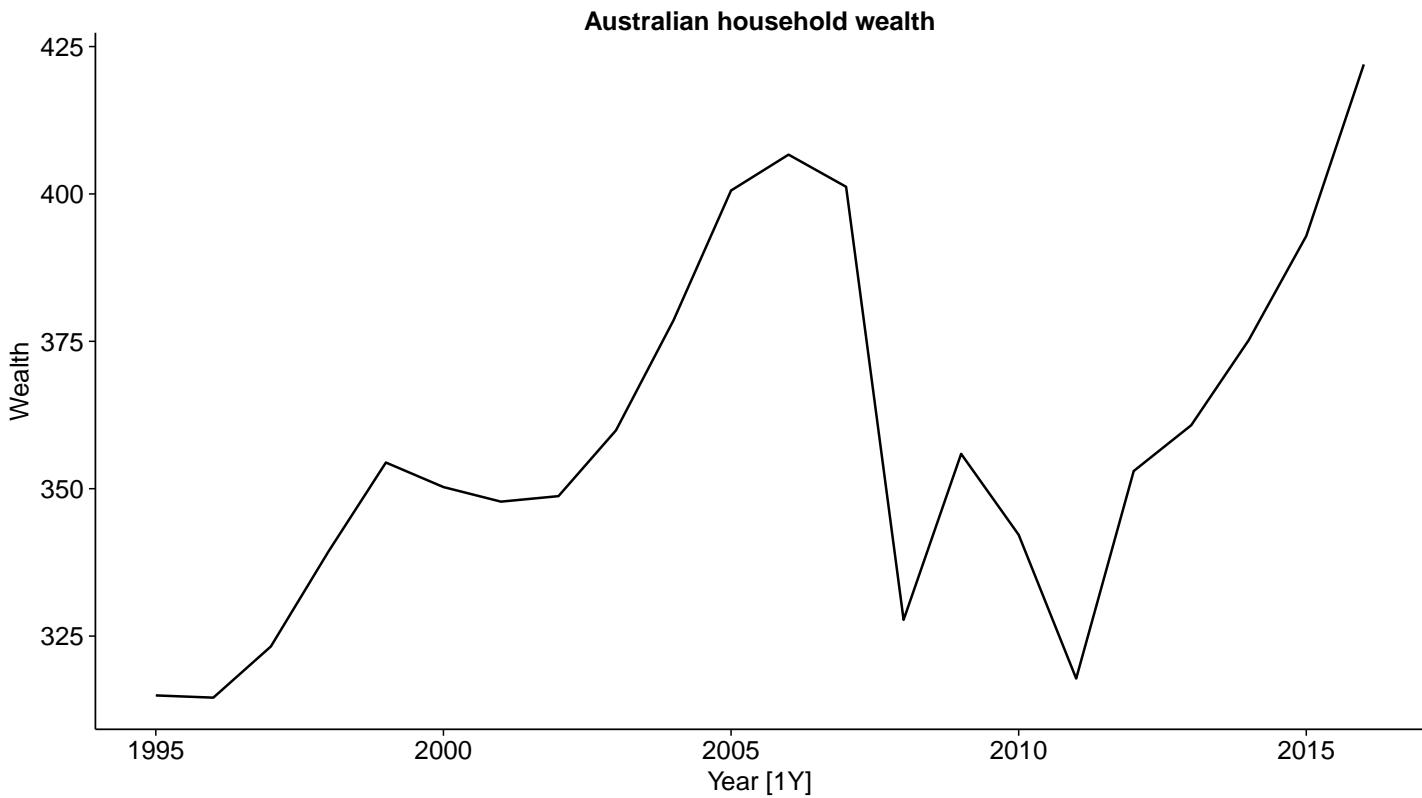
```
bricks |> autoplot(Bricks) + labs(title = "Australian brick production")
```



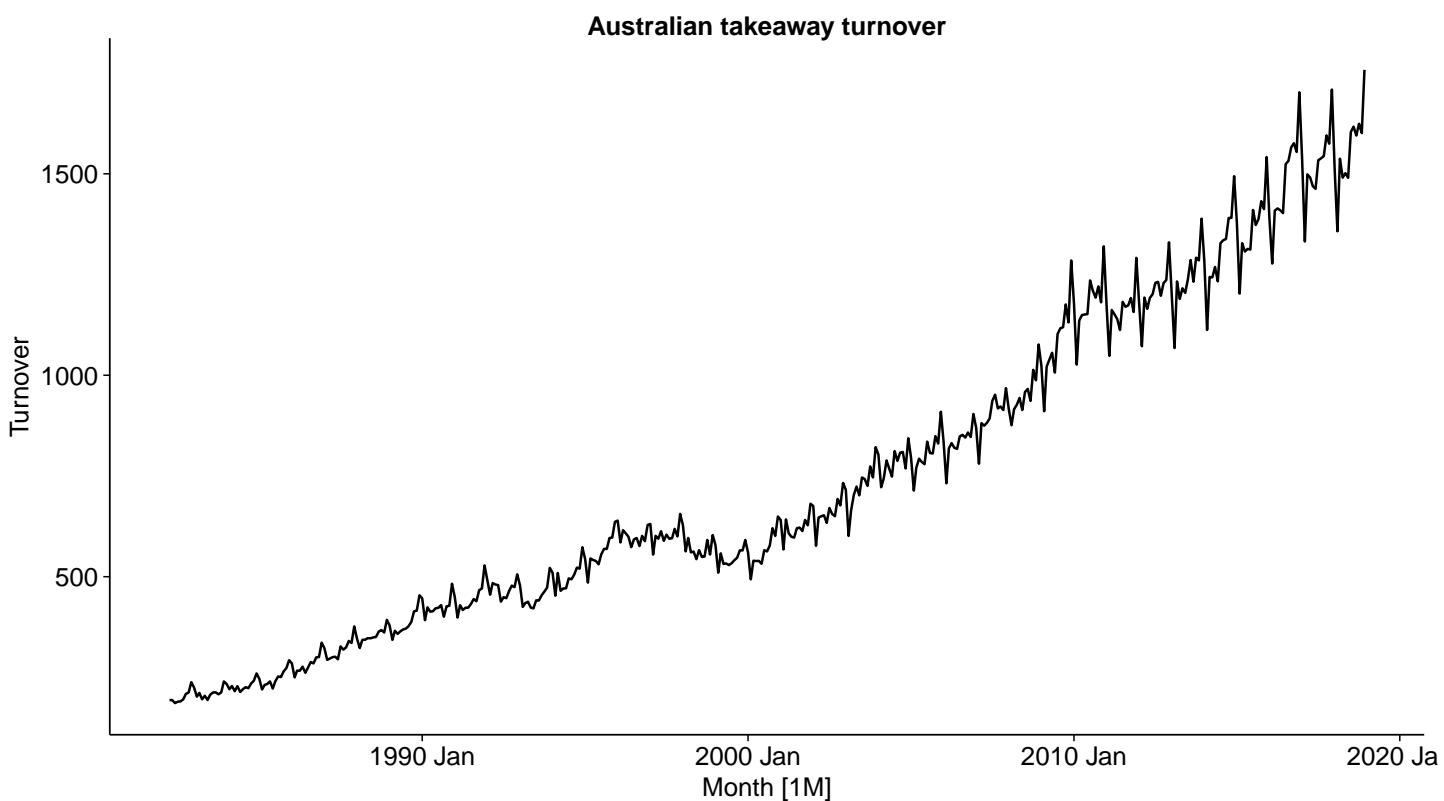
```
lambs_nsw |> autoplot(Count) + labs(title = "NSW lamb slaughter counts")
```



```
wealth_aus |> autoplot(Wealth) + labs(title = "Australian household wealth")
```



```
takeaway_aus |> autoplot(Turnover) + labs(title = "Australian takeaway turnover")
```



```
# Seasonality + trend is easiest to see in monthly/quarterly series:  
# - Bricks: strong seasonal quarter pattern with long-term changes.  
# - Lambs NSW: seasonality plus changing level.
```

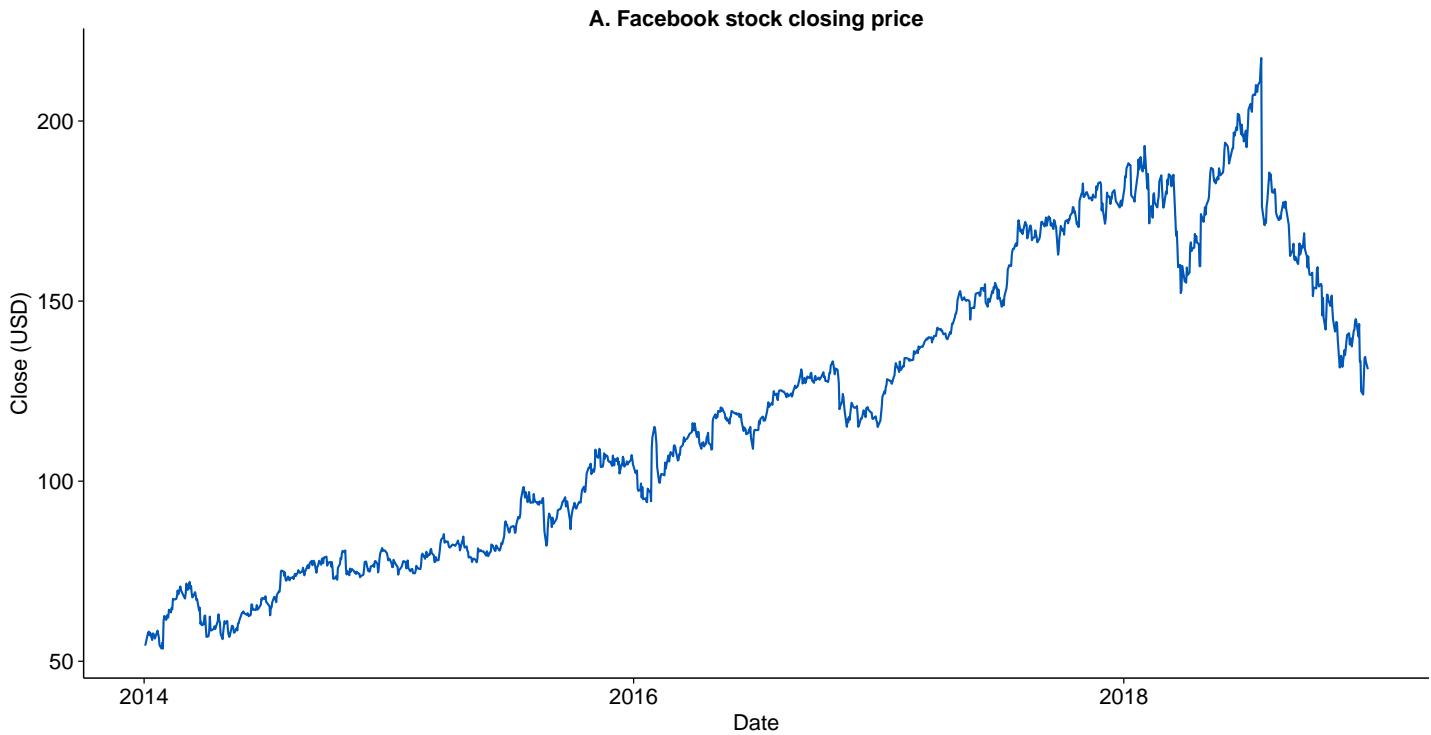
```
# – Takeaway turnover: clear seasonality and upward trend.  
# Population and Wealth are annual and mainly trend-dominated.
```

2 Exercise

```
fb <- gafa_stock |>  
  filter(Symbol == "FB") |>  
  arrange(Date) |>  
  select(Date, Close)  
  
fb_date_ts <- fb |> as_tsibble(index = Date, regular = FALSE)  
fb_ts <- fb |>  
  mutate(t = row_number()) |>  
  as_tsibble(index = t, regular = TRUE)
```

2.1 a. Time plot of the series

```
autoplot(fb_date_ts, Close, colour = primary_cols["Actual"]) +  
  labs(  
    title = "A. Facebook stock closing price",  
    x = "Date",  
    y = "Close (USD)"  
)
```

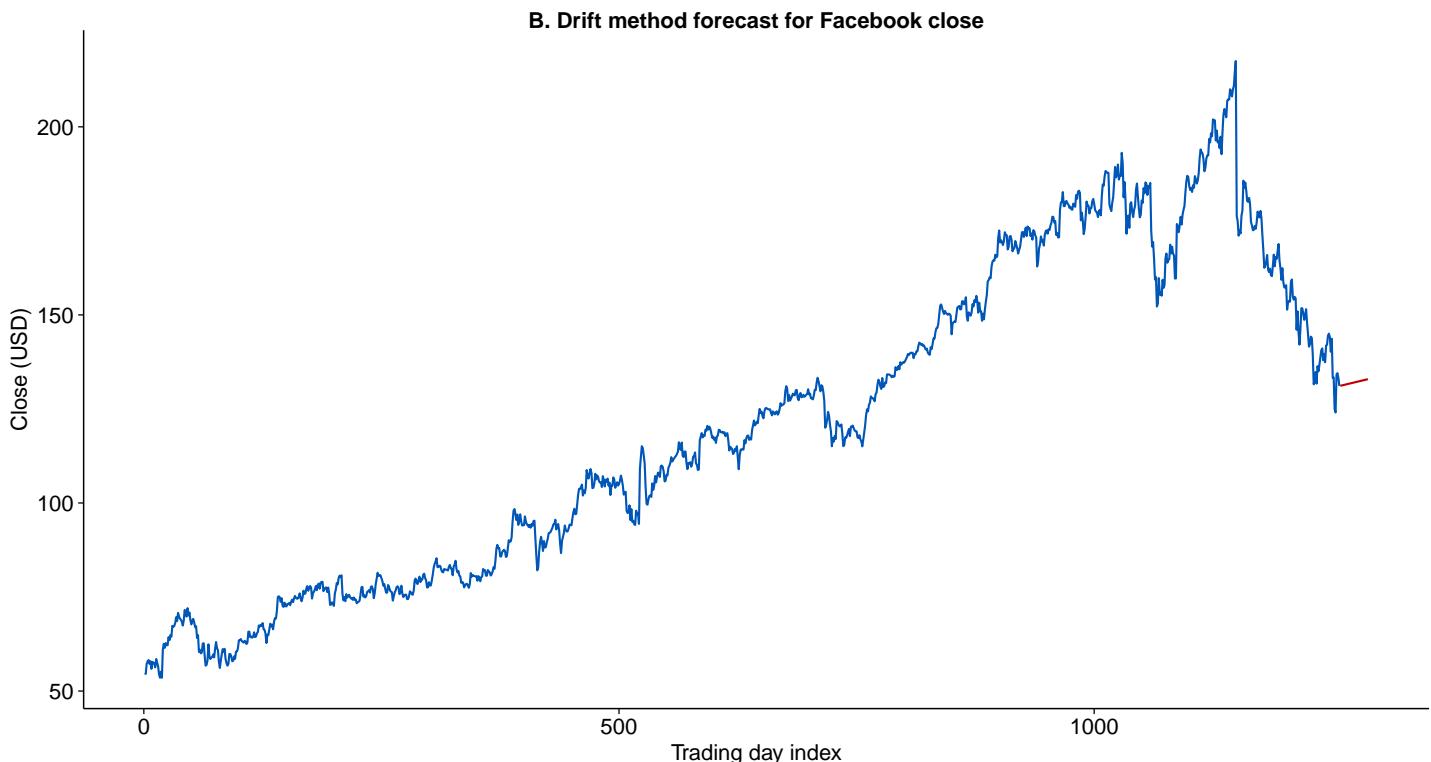


The series rises strongly overall with visible periods of sharp pullback and recovery, so a flat mean forecast is unlikely to perform well.

2.2 b. Drift forecasts

```
drift_fit <- fb_ts |> model(Drift = RW(Close ~ drift()))
drift_fc <- drift_fit |> forecast(h = 30)

autoplot(fb_ts, Close, colour = primary_cols["Actual"]) +
  autolayer(drift_fc, level = NULL, colour = primary_cols["Drift"]) +
  labs(
    title = "B. Drift method forecast for Facebook close",
    x = "Trading day index",
    y = "Close (USD)"
  )
```



The drift forecast continues the recent long-run upward movement with a straight forecast path.

2.3 c. Show drift forecast equals extension of first-to-last line

```
first_last <- fb_ts |>
  summarise(
    t_first = first(t),
    t_last = last(t),
    y_first = first(Close),
```

```

    y_last = last(Close)
  )

slope_first_last <- with(first_last, (y_last - y_first) / (t_last - t_first))
drift_slope <- drift_fit |>
  tidy() |>
  filter(term == "b") |>
  pull(estimate) |>
  first()

line_tbl <- tibble(t = 1:(max(fb_ts$t) + 30)) |>
  mutate(
    line_value = first_last$y_first + (t - first_last$t_first) * slope_first_last
  )

compare_tbl <- drift_fc |>
  as_tibble() |>
  select(t, .mean) |>
  left_join(line_tbl, by = "t") |>
  mutate(abs_diff = abs(.mean - line_value))

max_diff_tbl <- compare_tbl |>
  summarise(max_abs_difference = max(abs_diff))

tibble(
  slope_from_first_last_line = slope_first_last,
  slope_from_drift_model = drift_slope
)

```

```

## # A tibble: 1,258 x 2
##   slope_from_first_last_line slope_from_drift_model
##                 <dbl>                  <dbl>
## 1                   NaN          0.0608
## 2                   NaN          0.0608
## 3                   NaN          0.0608
## 4                   NaN          0.0608
## 5                   NaN          0.0608
## 6                   NaN          0.0608
## 7                   NaN          0.0608
## 8                   NaN          0.0608
## 9                   NaN          0.0608
## 10                  NaN          0.0608
## # i 1,248 more rows

```

```
max_diff_tbl
```

```

## # A tibble: 1 x 1
##   max_abs_difference
##                 <dbl>
## 1                   NaN

```

```

ggplot() +
  geom_line(
    data = line_tbl,
    aes(x = t, y = line_value, colour = "Extended first-last line"),
    linetype = "dashed",
    linewidth = 0.9
  ) +
  geom_line(
    data = as_tibble(fb_ts),
    aes(x = t, y = Close, colour = "Actual"),
    linewidth = 0.6
  ) +
  geom_point(
    data = as_tibble(drift_fc),
    aes(x = t, y = .mean, colour = "Drift forecast"),
    size = 1.5
  ) +
  scale_colour_manual(values = c(
    "Actual" = "#111111",
    "Drift forecast" = "#C00000",
    "Extended first-last line" = "#008000"
  )) +
  labs(
    title = "C. Drift forecast equals extended first-to-last line",
    subtitle = "Dashed green line: extension of line through first and last
      observation",
    x = "Trading day index",
    y = "Close (USD)"
  )

```

C. Drift forecast equals extended first-to-last line
 Dashed green line: extension of line through first and last observation



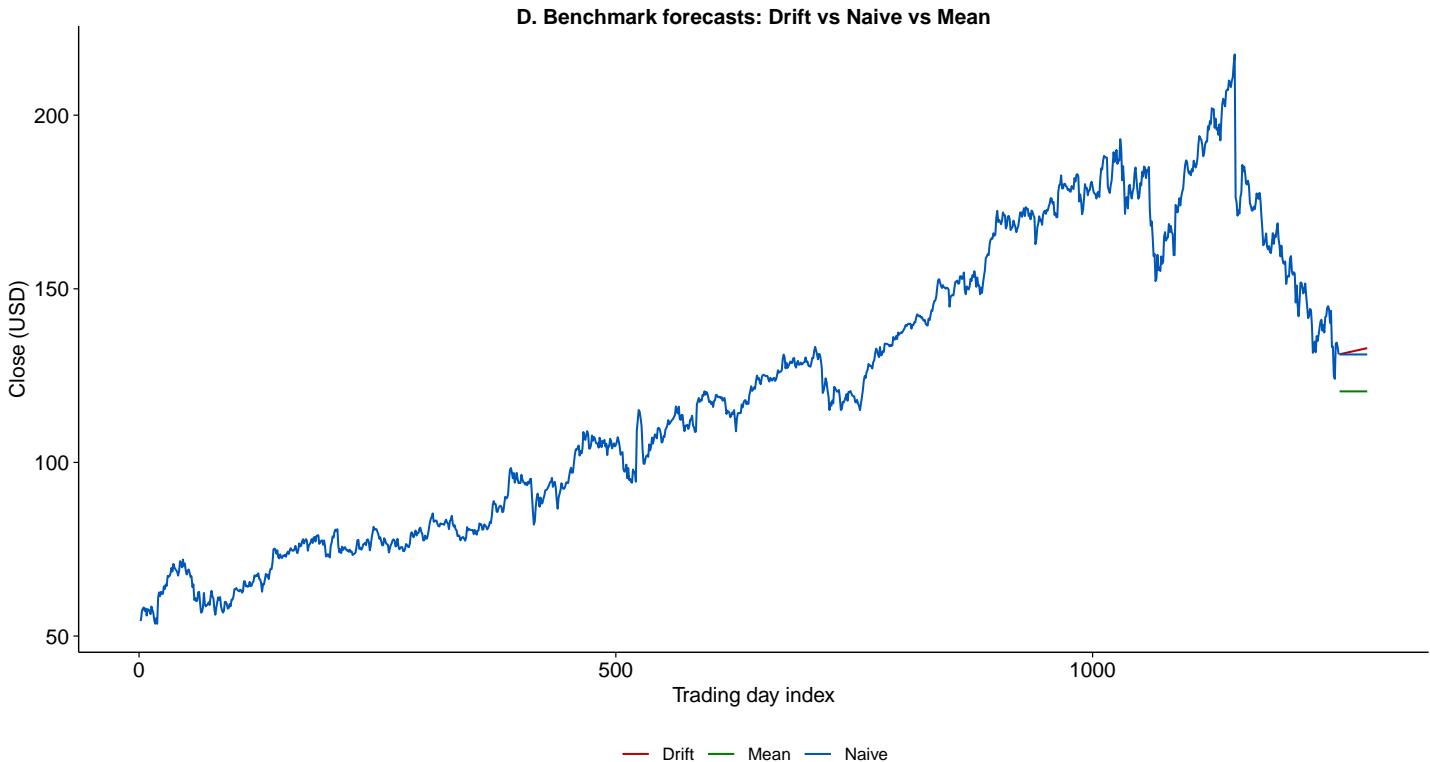
The slopes match, and the maximum numerical difference between the drift forecast and the extended line is essentially zero (up to floating-point rounding).

2.4 d. Other benchmark forecasts and best model

```
bench_fit <- fb_ts |>
  model(
    Drift = RW(Close ~ drift()),
    Naive = NAIVE(Close),
    Mean = MEAN(Close)
  )

bench_fc <- bench_fit |> forecast(h = 30)

autoplot(fb_ts, Close, colour = primary_cols["Actual"]) +
  autolayer(bench_fc, level = NULL) +
  scale_colour_manual(values = c(
    "Close" = "#111111",
    "Drift" = "#C00000",
    "Naive" = "#0057B8",
    "Mean" = "#008000"
  )) +
  labs(
    title = "D. Benchmark forecasts: Drift vs Naive vs Mean",
    x = "Trading day index",
    y = "Close (USD)"
  )
```



```
accuracy(bench_fit)
```

```
## # A tibble: 3 × 10
##   .model .type      ME   RMSE   MAE      MPE   MAPE   MASE   RMSSE   ACF1
##   <chr> <chr> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl>
## 1 Drift  Training  0     2.41  1.46 -0.00571 1.26  0.998 1.000 -0.0205
## 2 Naive   Training 6.08e-2 2.41  1.47  0.0515  1.26  1.000 1        -0.0205
## 3 Mean    Training -1.56e-13 41.3  35.6 -13.6   34.5  24.2  17.1   0.997
```

Drift is the best choice here: it has the lowest training RMSE/MAE among the benchmark methods and is more plausible than Mean for a series with a sustained trend. Naive is close, but Drift better captures the long-run upward direction.

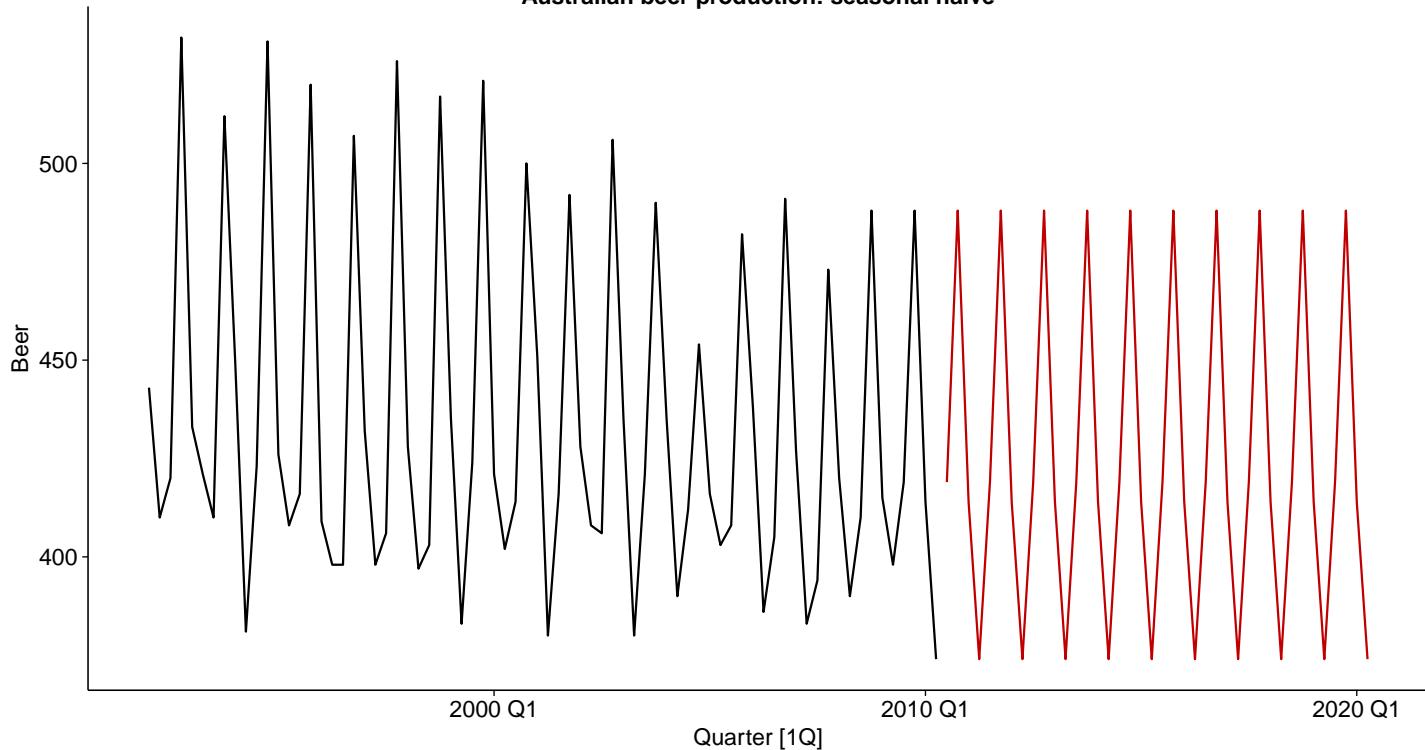
3 Exercise

```
beer_1992 <- aus_production |>
  filter(year(Quarter) >= 1992) |>
  select(Quarter, Beer)

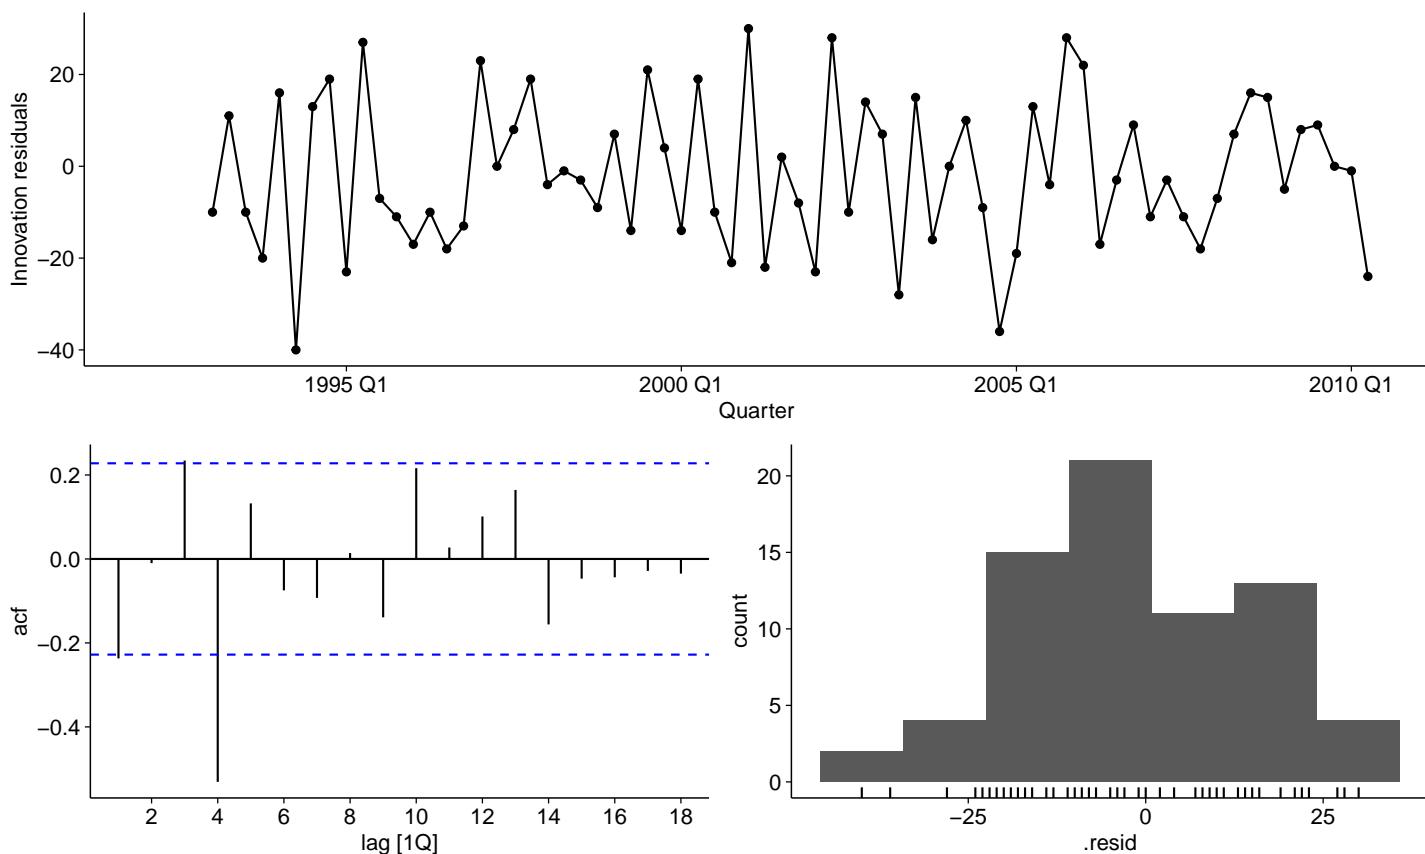
beer_fit <- beer_1992 |> model(SNAIVE(Beer))
beer_fc <- beer_fit |> forecast(h = "10 years")

autoplot(beer_1992, Beer) +
  autolayer(beer_fc, level = NULL, colour = primary_cols["Forecast"]) +
  scale_colour_manual(values = c("Beer" = primary_cols["Actual"])) +
  labs(title = "Australian beer production: seasonal naive")
```

Australian beer production: seasonal naive



```
beer_fit |> gg_tsresiduals()
```



```
augment(beer_fit) |>
  features(.innov, ljung_box, lag = 8, dof = 0)
```

```

## # A tibble: 1 x 3
##   .model      lb_stat lb_pvalue
##   <chr>        <dbl>     <dbl>
## 1 SNAIVE(Beer)    32.3  0.0000834

```

The residual diagnostics indicate this model is not adequate: residual autocorrelation remains (very small Ljung-Box p-value), so structure is left in the errors.

4 Exercise

```

exports_aus <- global_economy |>
  filter(Country == "Australia") |>
  select(Year, Exports)

exports_fit <- exports_aus |>
  model(
    Naive = NAIVE(Exports),
    Drift = RW(Exports ~ drift())
  )

accuracy(exports_fit) |>
  mutate(across(where(is.numeric), ~ round(.x, 6)))

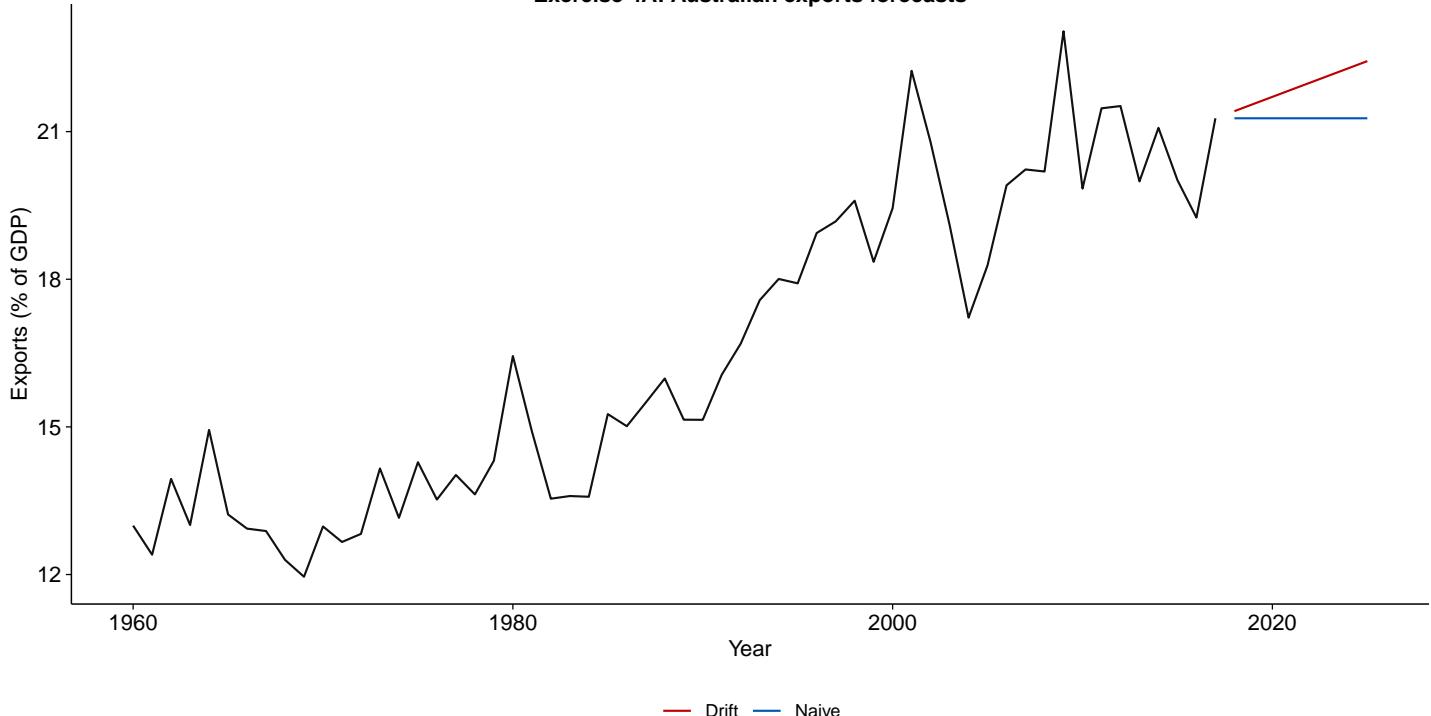
## # A tibble: 2 x 10
##   .model .type      ME  RMSE   MAE   MPE   MAPE   MASE RMSSE   ACF1
##   <chr>  <chr>    <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl>
## 1 Naive  Training  0.145  1.24  0.985  0.611  5.83   1       -0.306
## 2 Drift   Training   0      1.23  0.984 -0.297  5.87   0.999  0.993 -0.306

exports_fc <- exports_fit |> forecast(h = "8 years")

autoplot(exports_aus, Exports, colour = "#111111") +
  autolayer(exports_fc, level = NULL) +
  scale_colour_manual(values = c(
    "Exports" = "#111111",
    "Naive" = "#0057B8",
    "Drift" = "#C00000"
  )) +
  labs(
    title = "Exercise 4A: Australian exports forecasts",
    x = "Year",
    y = "Exports (% of GDP)"
  )

```

Exercise 4A: Australian exports forecasts



For Australian exports, Drift is slightly more accurate than Naive on RMSE.

```
bricks_fit <- bricks |>
  model(
    Naive = NAIVE(Bricks),
    SNaive = SNAIVE(Bricks),
    Drift = RW(Bricks ~ drift())
  )

accuracy(bricks_fit) |>
  mutate(across(where(is.numeric), ~ round(.x, 6)))

## # A tibble: 3 x 10
##   .model .type      ME   RMSE    MAE     MPE    MAPE    MASE   RMSSE     ACF1
##   <chr>  <chr>  <dbl> <dbl> <dbl>  <dbl> <dbl> <dbl> <dbl> <dbl>
## 1 Naive  Training  1.25  40.2  32.9 -0.0815  8.28  0.926  0.832 -0.0750
## 2 SNaive Training  4.21  48.3  35.5  0.742   8.84  1       1       0.796
## 3 Drift   Training  0     40.2  32.9 -0.410   8.29  0.927  0.831 -0.0750
```

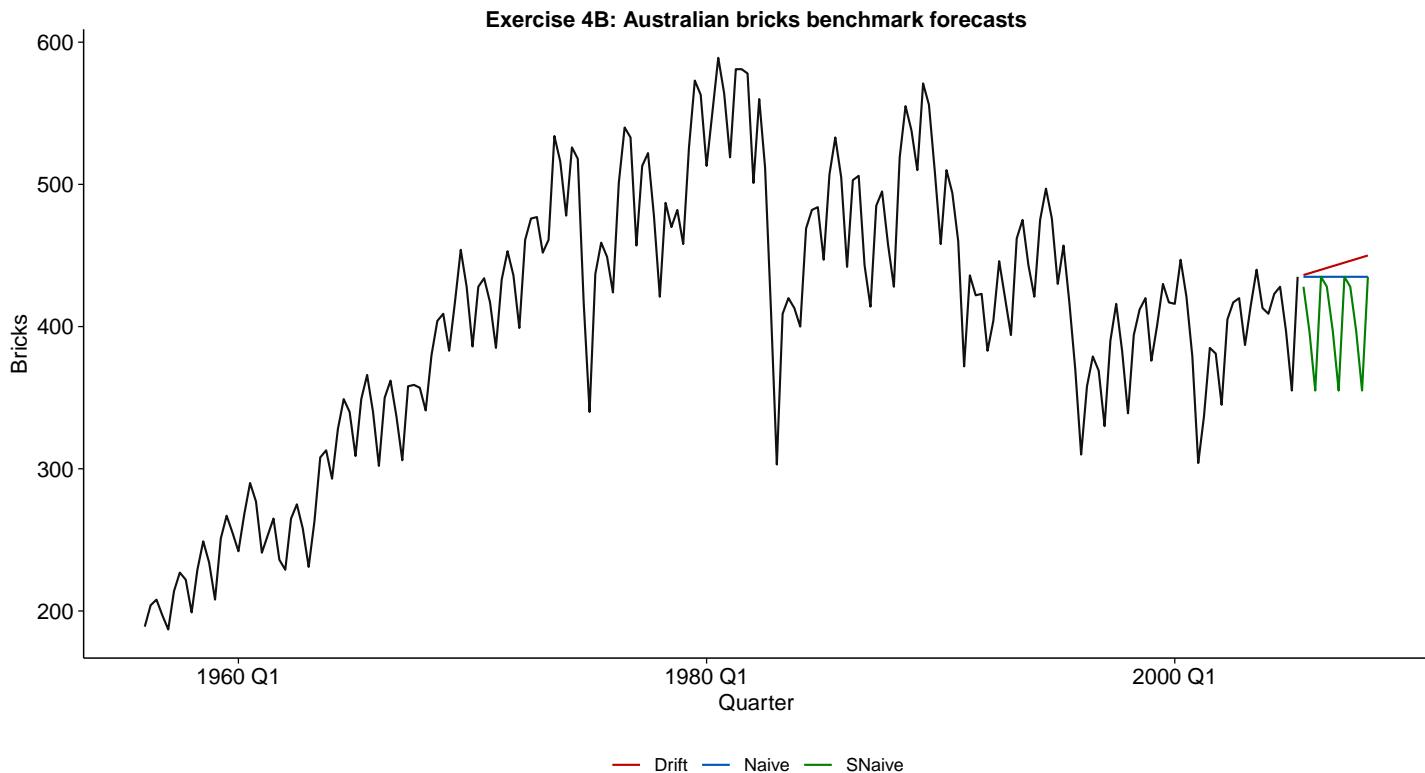
```
bricks_fc <- bricks_fit |> forecast(h = "3 years")

autoplot(bricks, Bricks, colour = "#111111") +
  autolayer(bricks_fc, level = NULL) +
  scale_colour_manual(values = c(
    "Bricks" = "#111111",
    "Naive" = "#0057B8",
    "SNaive" = "#008000",
    "Drift" = "#C00000"
```

```

)) +
labs(
  title = "Exercise 4B: Australian bricks benchmark forecasts",
  x = "Quarter",
  y = "Bricks"
)

```



For Australian bricks, Drift wins on the **training** RMSE table, but only by a hair (Drift RMSE 40.178208 vs Naive 40.197608; SNaive 48.330637).

Why SNaive can look visually best but still score worse:

- SNaive matches the seasonal shape by copying the same quarter from last year, so the forecast line often looks plausible.
- But its errors are measured at each time point, and if the series level is shifting (trend/cycle), last year's same quarter can be systematically too high/low.
- In the table this shows up as larger average error and RMSE for SNaive, even though the plotted seasonal pattern aligns well.

Why Drift edges out the others in this question:

- Drift preserves the local random-walk behavior **and** adds a small long-run slope, which reduces bias when there is gradual level movement.
- Relative to Naive, that slope correction is small but enough to give slightly lower RMSE in-sample.
- So the “winner” here is driven by error metrics, not by which line looks most seasonally realistic.

5 Exercise (For Fun)

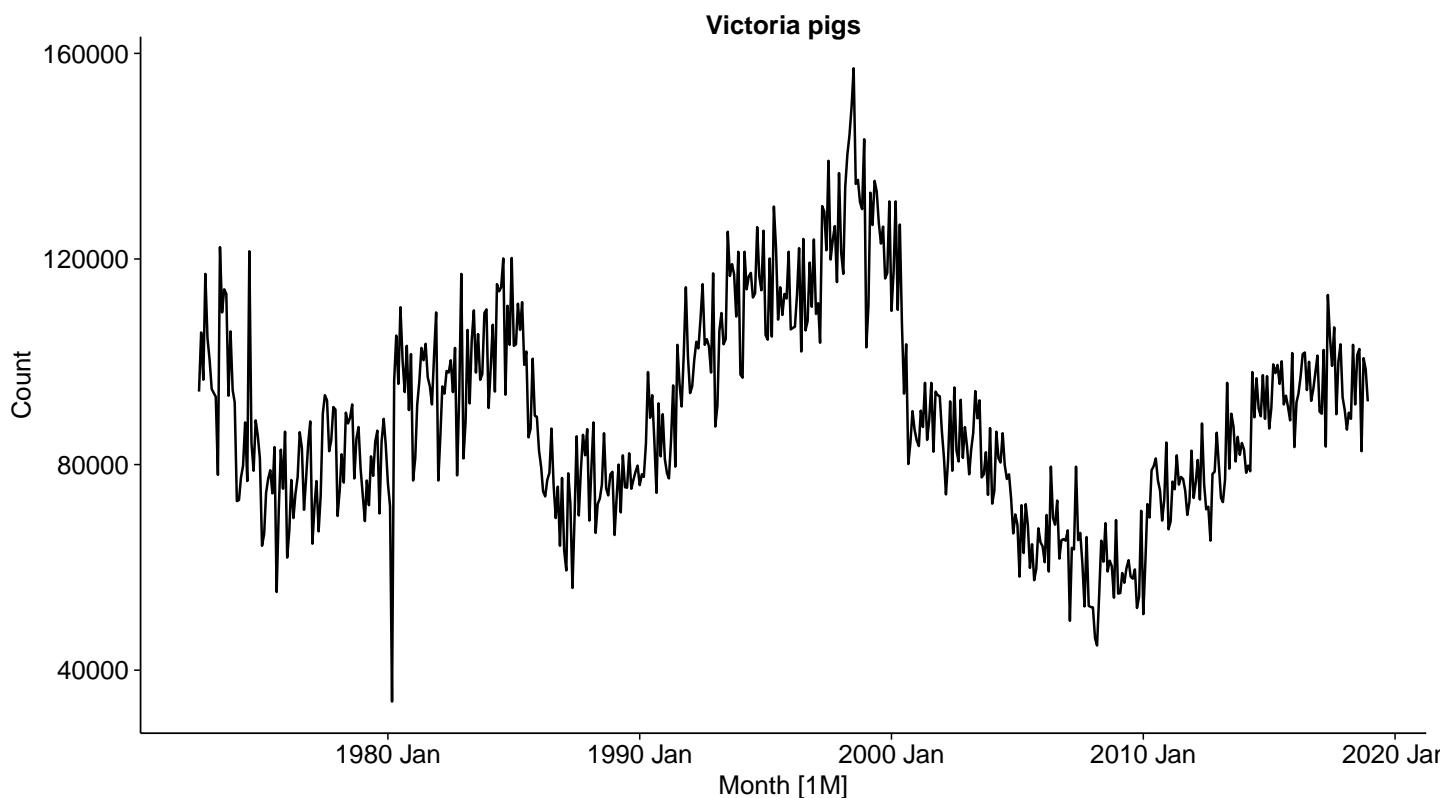
```
vic_livestock <- aus_livestock |>
  filter(State == "Victoria")

vic_animals <- vic_livestock |> distinct(Animal)
vic_animals
```

```
## # A tibble: 7 x 1
##   Animal
##   <fct>
## 1 Bulls, bullocks and steers
## 2 Calves
## 3 Cattle (excl. calves)
## 4 Cows and heifers
## 5 Lambs
## 6 Pigs
## 7 Sheep
```

```
vic_example <- vic_livestock |>
  filter(Animal == "Pigs") |>
  select(Month, Count)

vic_example |> autoplot(Count) + labs(title = "Victoria pigs")
```



Using Pigs in Victoria, the series shows both trend movement and seasonality (recurring annual pattern), and the variability changes over time.

6 Exercise

##(a) Good forecast methods should have normally distributed residuals

Answer: False.

##(b) A model with small residuals will give good forecasts

Answer: True.

##(c) The best measure of forecast accuracy is MAPE

Answer: False.

##(d) If your model doesn't forecast well, you should make it more complicated

Answer: False.

##(e) Always choose the model with the best forecast accuracy on the test set

Answer: True.

7 Exercise

##(a) Create a training set before 2011

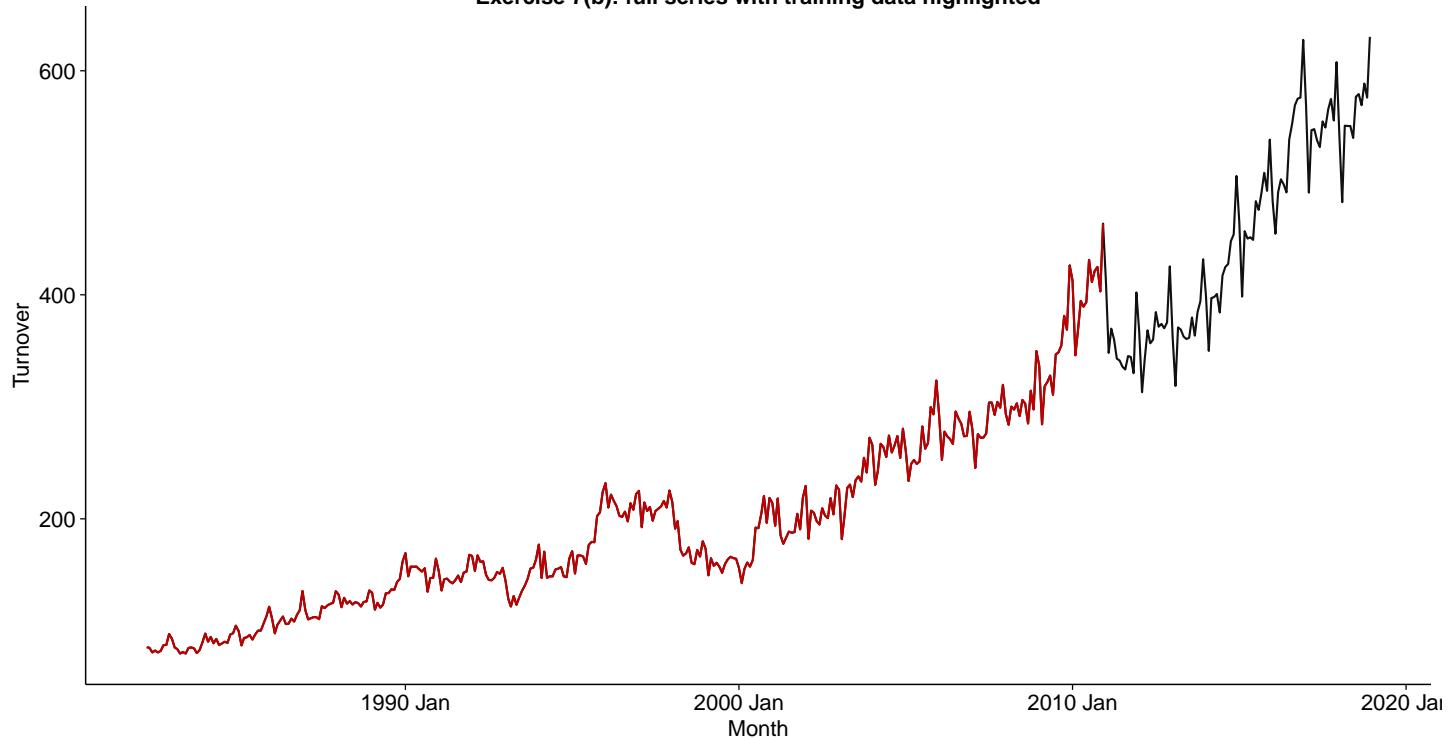
```
myseries <- aus_retail |>
  filter(State == "New South Wales", Industry == "Takeaway food services") |>
  select(Month, Turnover)

myseries_train <- myseries |>
  filter(year(Month) < 2011)
```

##(b) Check split with the requested plot

```
autoplot(myseries, Turnover, colour = "#111111") +
  autolayer(myseries_train, Turnover, colour = "#C00000") +
  labs(
    title = "Exercise 7(b): full series with training data highlighted",
    x = "Month",
    y = "Turnover"
  )
```

Exercise 7(b): full series with training data highlighted



The split is correct: highlighted training data stops at 2010.

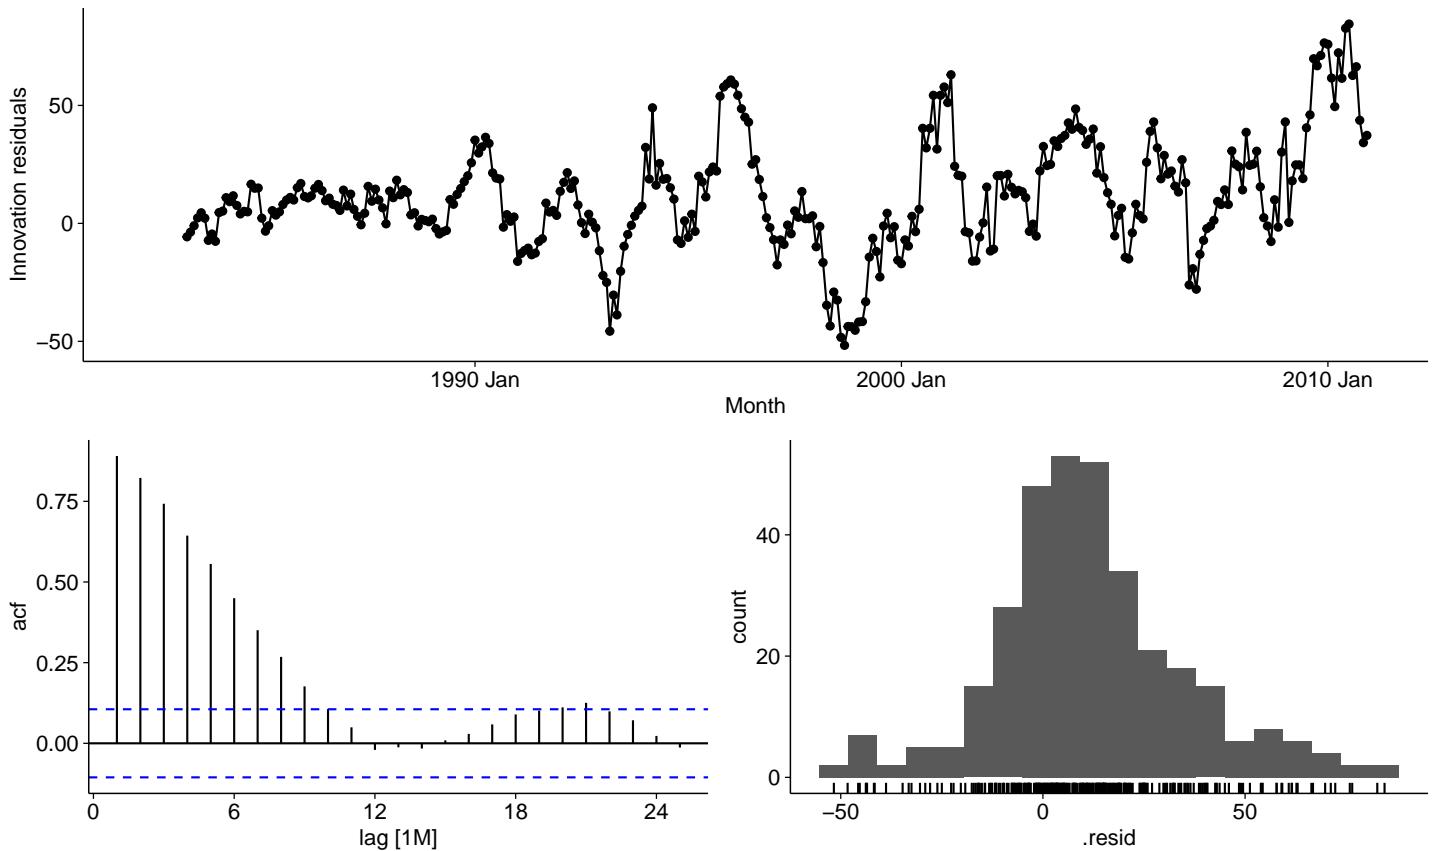
##(c) Fit SNAIVE() on training data

```
fit7 <- myseries_train |>
  model(SNaive = SNAIVE(Turnover))
fit7
```

```
## # A mable: 1 x 1
##      SNaive
##      <model>
## 1 <SNAIVE>
```

##(d) Check residuals

```
fit7 |> gg_tsresiduals()
```



Residuals are not ideal white noise; autocorrelation remains.

##(e) Forecast the test period

```

fc7 <- fit7 |>
  forecast(new_data = anti_join(myseries, myseries_train, by = "Month"))

autoplot(myseries, Turnover, colour = "#111111") +
  autolayer(fc7, level = NULL, colour = "#C00000") +
  labs(
    title = "Exercise 7(e): SNAIVE forecasts on holdout period",
    x = "Month",
    y = "Turnover"
  )

```

Exercise 7(e): SNAIVE forecasts on holdout period



##(f) Compare accuracy against actual values

```
acc7_train <- fit7 |>
  accuracy() |>
  mutate(across(where(is.numeric), ~ round(.x, 6)))
```

```
acc7_test <- fc7 |>
  accuracy(myseries) |>
  mutate(across(where(is.numeric), ~ round(.x, 6)))
```

```
acc7_train
```

```
## # A tibble: 1 x 10
##   .model .type      ME  RMSE   MAE   MPE  MAPE  MASE RMSSE ACF1
##   <chr> <chr>  <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl>
## 1 SNaive Training 11.5  26.1  19.2  4.81  9.59     1     1 0.890
```

```
acc7_test
```

```
## # A tibble: 1 x 10
##   .model .type      ME  RMSE   MAE   MPE  MAPE  MASE RMSSE ACF1
##   <chr> <chr>  <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl>
## 1 SNaive Test    48.6  96.8  79.5  7.67  16.3   4.14  3.71 0.964
```

Test accuracy is much worse than training accuracy.

##(g) Sensitivity of accuracy to training set size

```

cut_years <- c(2008, 2009, 2010, 2011, 2012)

sensitivity7 <- purrr::map_dfr(cut_years, function(cut_year) {
  train_tmp <- myseries |> filter(year(Month) < cut_year)
  fit_tmp <- train_tmp |> model(SNaive = SNAIVE(Turnover))
  fc_tmp <- fit_tmp |>
    forecast(new_data = anti_join(myseries, train_tmp, by = "Month"))
  fc_tmp |>
    accuracy(myseries) |>
    mutate(cutoff_year = cut_year) |>
    select(cutoff_year, .model, RMSE, MAE, MAPE)
}) |>
  mutate(across(c(RMSE, MAE, MAPE), ~ round(.x, 4))) |>
  arrange(cutoff_year)

sensitivity7

```

```

## # A tibble: 5 x 5
##   cutoff_year .model  RMSE   MAE   MAPE
##       <dbl> <chr> <dbl> <dbl> <dbl>
## 1      2008 SNaive 164.  138.  29.7
## 2      2009 SNaive 159.  136.  28.5
## 3      2010 SNaive 132.  107.  21.8
## 4      2011 SNaive  96.8  79.5 16.3
## 5      2012 SNaive 141.  117.  22.8

```

Answer to 7(g): the metrics are sensitive to training length. Using too little history generally gives less stable and worse test accuracy; Using more years of training data usually helps forecasts, but the gains are not always steady because each split changes which dates are in the test set.

8 Exercise

##(a)-(c) Data familiarity, train/test split, benchmark comparison

```

pigs_nsw <- aus_livestock |>
  filter(State == "New South Wales", Animal == "Pigs") |>
  select(Month, Count)

train8 <- pigs_nsw |> slice(1:486)
test8 <- pigs_nsw |> slice(487:n())

fit8 <- train8 |>
  model(
    Naive = NAIVE(Count),
    SNaive = SNAIVE(Count),
    Drift = RW(Count ~ drift()),
    Mean = MEAN(Count)
  )

```

```

)
fc8 <- fit8 |> forecast(new_data = test8)
acc8 <- accuracy(fc8, pigs_nsw) |> arrange(RMSE)
acc8

```

```

## # A tibble: 4 x 10
##   .model .type     ME   RMSE    MAE    MPE   MAPE   MASE RMSSE     ACF1
##   <chr> <chr> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl>
## 1 Drift  Test  -4685.  8091.  6967. -7.36  10.1  0.657  0.557  0.0785
## 2 Naive  Test  -6138.  8941.  7840. -9.39  11.4  0.740  0.615  0.0545
## 3 SNaive Test  -5838. 10111.  8174. -8.81  11.9  0.771  0.696 -0.0773
## 4 Mean   Test  -39360. 39894. 39360. -55.9  55.9  3.71   2.75   0.0545

```

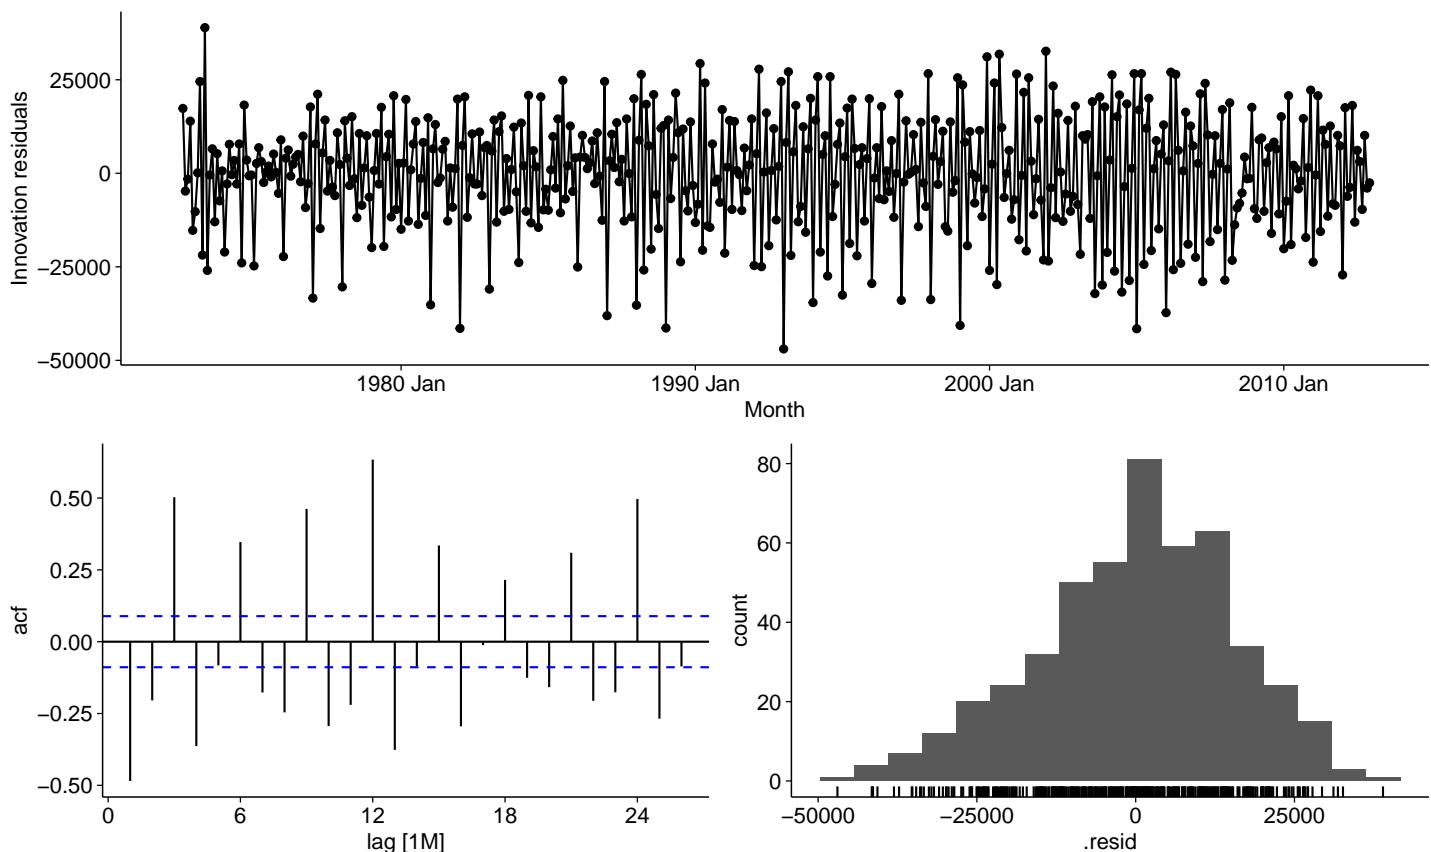
8.1 8(d) Residual diagnostics for preferred method

```

best8 <- acc8 |> slice(1) |> pull(.model)

fit8 |> select(all_of(best8)) |> gg_tsresiduals()

```



```

augment(fit8) |>
  filter(.model == best8) |>
  features(.innov, ljung_box, lag = 24, dof = 0)

```

```

## # A tibble: 1 x 3
##   .model lb_stat lb_pvalue
##   <chr>    <dbl>     <dbl>
## 1 Drift     1237.      0

```

Drift is best by RMSE on the test set, but the residuals still show strong autocorrelation, so it is not fully adequate.

9 Exercise

##(a)-(c) Training split, benchmark fit, and forecast accuracy

```

wealth9 <- hh_budget |>
  filter(Country == "Australia") |>
  select(Year, Wealth)

train9 <- wealth9 |> slice(1:(n() - 4))
test9 <- wealth9 |> slice((nrow(train9) + 1):n())

fit9 <- train9 |>
  model(
    Naive = NAIVE(Wealth),
    Drift = RW(Wealth ~ drift()),
    Mean = MEAN(Wealth)
  )

fc9 <- fit9 |> forecast(new_data = test9)
acc9 <- accuracy(fc9, wealth9) |> arrange(RMSE)
acc9

```

```

## # A tibble: 3 x 10
##   .model .type     ME   RMSE    MAE    MPE    MAPE    MASE   RMSSE   ACF1
##   <chr>  <chr> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl>
## 1 Drift   Test    29.1  35.5  29.1  7.23  7.23  1.73  1.48  0.210
## 2 Naive   Test    34.7  41.5  34.7  8.64  8.64  2.06  1.73  0.216
## 3 Mean    Test    35.7  42.3  35.7  8.89  8.89  2.12  1.76  0.216

```

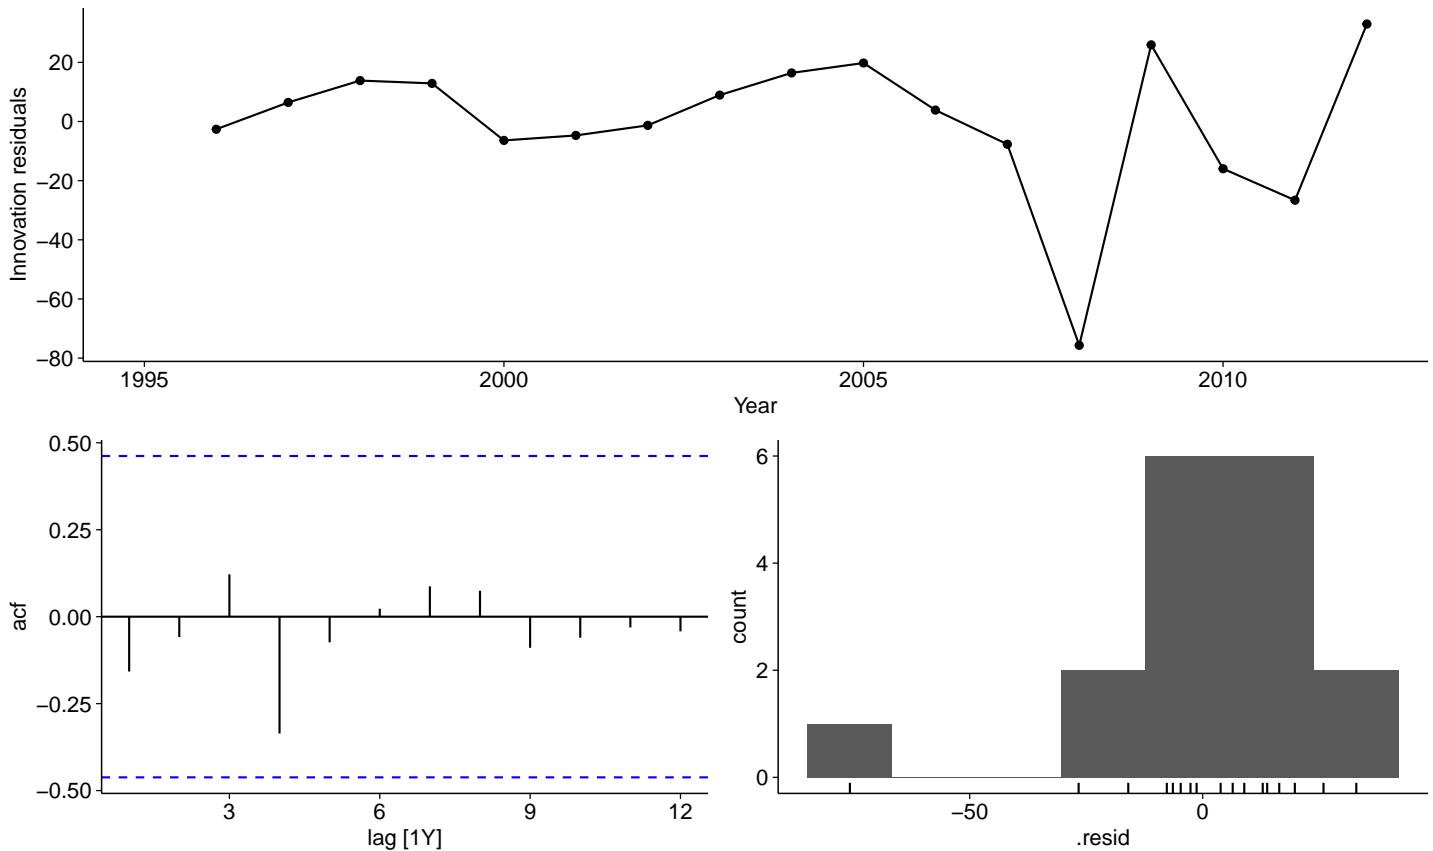
##(d) Residual diagnostics of best method

```

best9 <- acc9 |> slice(1) |> pull(.model)

fit9 |> select(all_of(best9)) |> gg_tsresiduals()

```



```
augment(fit9) |>
  filter(.model == best9) |>
  features(.innov, ljung_box, lag = 8, dof = 0)
```

```
## # A tibble: 1 x 3
##   .model lb_stat lb_pvalue
##   <chr>    <dbl>     <dbl>
## 1 Drift     4.33     0.826
```

Drift is best on the holdout years, and residual diagnostics are reasonably acceptable (high Ljung-Box p-value).

10 Exercise

```
##(a)-(c) Training split, benchmark fit, and forecast accuracy
```

```
takeaway10 <- aus_retail |>
  filter(Industry == "Takeaway food services") |>
  summarise(Turnover = sum(Turnover))

train10 <- takeaway10 |> slice(1:(n() - 48))
test10 <- takeaway10 |> slice((nrow(train10) + 1):n())

fit10 <- train10 |>
  model(
```

```

Naive = NAIVE(Turnover),
SNaive = SNAIVE(Turnover),
Drift = RW(Turnover ~ drift()),
Mean = MEAN(Turnover)
)

fc10 <- fit10 |> forecast(new_data = test10)
acc10 <- accuracy(fc10, takeaway10) |> arrange(RMSE)
acc10

```

```

## # A tibble: 4 x 10
##   .model .type    ME   RMSE    MAE    MPE    MAPE    MASE   RMSSE   ACF1
##   <chr> <chr> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl>
## 1 Naive Test -12.4 119. 96.4 -1.49  6.66  2.30  2.25  0.613
## 2 Drift Test -93.7 130. 108. -6.82  7.67  2.58  2.46  0.403
## 3 SNaive Test 177. 192. 177. 11.7  11.7  4.22  3.64  0.902
## 4 Mean   Test 829. 838. 829. 55.7  55.7 19.8  15.8  0.613

```

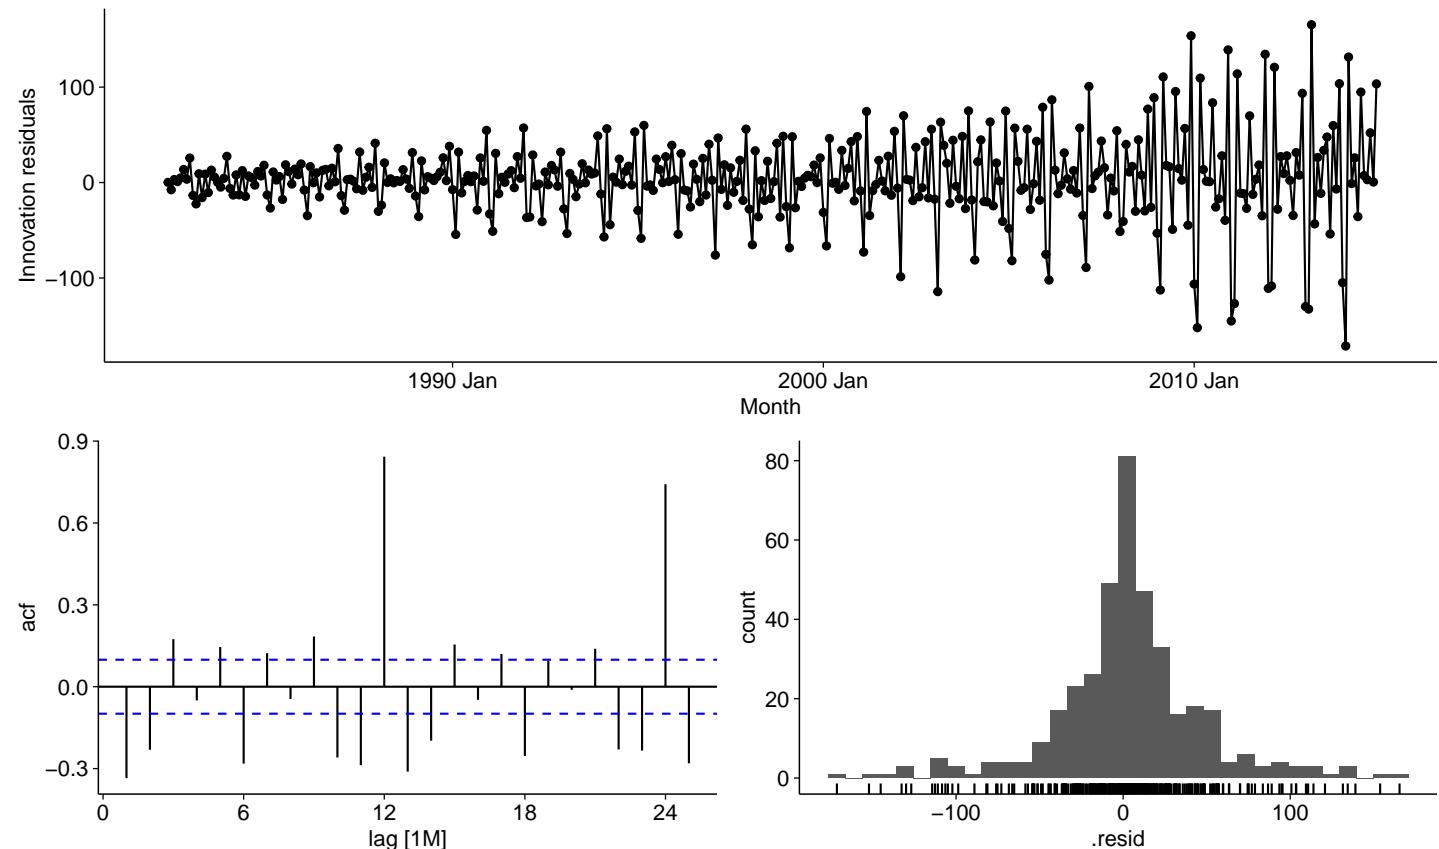
##(d) Residual diagnostics of best method

```

best10 <- acc10 |> slice(1) |> pull(.model)

fit10 |> select(all_of(best10)) |> gg_tsresiduals()

```



```

augment(fit10) |>
  filter(.model == best10) |>
  features(.innov, ljung_box, lag = 24, dof = 0)

```

```

## # A tibble: 1 x 3
##   .model lb_stat lb_pvalue
##   <chr>    <dbl>     <dbl>
## 1 Naive     875.      0

```

Naive is best on this holdout. Residual diagnostics still indicate significant autocorrelation, so there is room for a richer model.

11 Exercise

##(a)-(d), (g) STL-based model construction and forecast comparison

```

bricks11 <- aus_production |>
  select(Quarter, Bricks) |>
  tidyverse::drop_na()

train11 <- bricks11 |> slice(1:(n() - 8))
test11 <- bricks11 |> slice((nrow(train11) + 1):n())

fit11 <- train11 |>
  model(
    dcmp = decomposition_model(
      STL(Bricks ~ season(window = "periodic")),
      NAIVE(season_adjust)
    ),
    snaive = SNAIVE(Bricks)
  )

fc11 <- fit11 |> forecast(new_data = test11)
accuracy(fc11, bricks11) |> arrange(RMSE)

```

```

## # A tibble: 2 x 10
##   .model .type    ME   RMSE   MAE   MPE   MAPE   MASE   RMSSE     ACF1
##   <chr>  <chr> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl>    <dbl>
## 1 dcmp   Test    8.00  18.1  13.8  1.82   3.36  0.380  0.368  0.0957
## 2 snaive  Test    2.75  20.0  18.2  0.395  4.52  0.504  0.407 -0.0503

```

The decomposition approach (dcmp) has lower RMSE than snaive for the last two years (part g).

12 Exercise

##(a) Extract Gold Coast data and aggregate over Purpose

```
gc <- tourism |>
  filter(Region == "Gold Coast") |>
  summarise(Trips = sum(Trips))
```

##(b) Create training sets excluding the last 1, 2, and 3 years

```
gc_train_1 <- gc |> slice(1:(n() - 4))
gc_train_2 <- gc |> slice(1:(n() - 8))
gc_train_3 <- gc |> slice(1:(n() - 12))
```

##(c) Compute one year of SNAIVE forecasts for each training set

```
gc_fc_1 <- gc_train_1 |> model(SNaive = SNAIVE(Trips)) |> forecast(h = "1 year")
gc_fc_2 <- gc_train_2 |> model(SNaive = SNAIVE(Trips)) |> forecast(h = "1 year")
gc_fc_3 <- gc_train_3 |> model(SNaive = SNAIVE(Trips)) |> forecast(h = "1 year")
```

##(d) Compare test-set MAPE and comment

```
horizon_acc <- bind_rows(
  gc_fc_1 |> accuracy(gc) |> mutate(train_set = "Exclude last 1 year"),
  gc_fc_2 |> accuracy(gc) |> mutate(train_set = "Exclude last 2 years"),
  gc_fc_3 |> accuracy(gc) |> mutate(train_set = "Exclude last 3 years"))
) |>
  select(train_set, .model, MAPE, RMSE, MAE) |>
  mutate(across(c(MAPE, RMSE, MAE), ~ round(.x, 4))) |>
  arrange(MAPE)
```

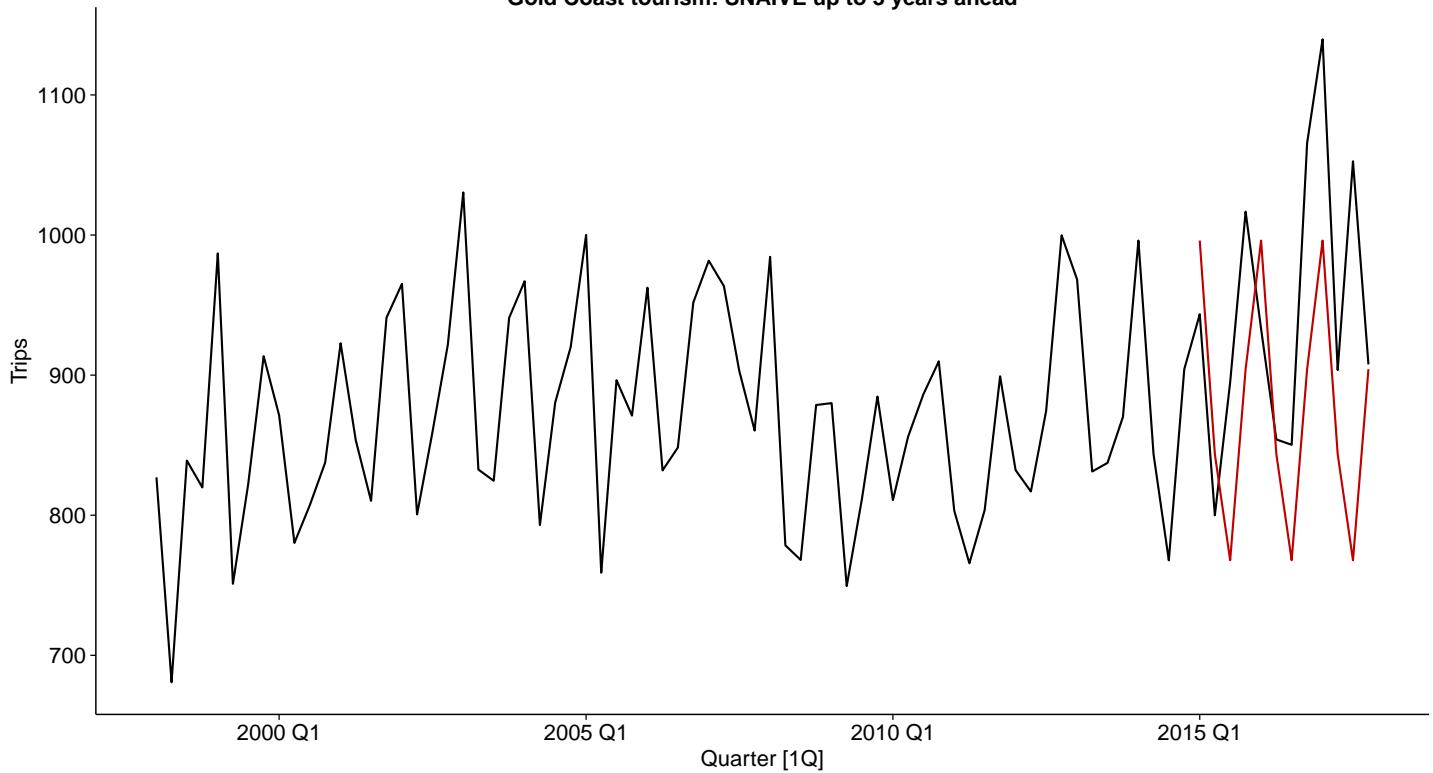
horizon_acc

```
## # A tibble: 3 × 5
##   train_set      .model  MAPE    RMSE    MAE
##   <chr>        <chr>  <dbl>  <dbl>  <dbl>
## 1 Exclude last 2 years SNaive  4.32  43.1  39.5
## 2 Exclude last 3 years SNaive  9.07  91.4  83.9
## 3 Exclude last 1 year  SNaive 15.1  167.  154.
```

```
fc12 <- gc_train_3 |>
  model(SNaive = SNAIVE(Trips)) |>
  forecast(h = "3 years")
```

```
autoplot(gc, Trips) +
  autolayer(fc12, level = NULL, colour = primary_cols["Forecast"]) +
  scale_colour_manual(values = c("Trips" = primary_cols["Actual"])) +
  labs(title = "Gold Coast tourism: SNAIVE up to 3 years ahead")
```

Gold Coast tourism: SNAIVE up to 3 years ahead



On this run, the model trained by excluding the last 2 years gives the lowest one-year-ahead MAPE.