

FPP3 Exercises Homework 6

Randy Howk

2026-04-11

Contents

1	Exercise 1. Pigs in Victoria	3
1.1	Exercise 1(a) Use ETS () to estimate the equivalent SES model, report α and λ_0 , and forecast four months ahead.	3
1.2	Exercise 1(b) Compute a 95% prediction interval for the first forecast using $\hat{y}$ $\pm 1.96 s$, and compare it with R's interval.	4
2	Exercise 2. Write your own SES function	5
3	Exercise 3. Modify the function to return SSE and optimize with <code>optim()</code>	6
4	Exercise 4. One function that estimates and forecasts	7
5	Exercise 5. Exports from one country in <code>global_economy</code>	7
5.1	Exercise 5(a) Plot the series and discuss the main features.	8
5.2	Exercise 5(b) Forecast with an ETS (A,N,N) model and plot the forecasts.	8
5.3	Exercise 5(c) Compute the training-set RMSE.	9
5.4	Exercise 5(d) Compare with ETS (A,A,N) and discuss the merits of the two methods.	9
5.5	Exercise 5(e) Compare the forecasts. Which is best?	10
5.6	Exercise 5(f) Compute a 95% prediction interval for the first forecast for each model using RMSE and compare with R.	11
6	Exercise 6. Chinese GDP under different ETS options	11
7	Exercise 7. Gas from <code>aus_production</code>	13
8	Exercise 8. Retail series revisited	14
8.1	Exercise 8(a) Why is multiplicative seasonality necessary?	14
8.2	Exercise 8(b) Apply Holt-Winters multiplicative and experiment with damping.	15
8.3	Exercise 8(c) Compare one-step RMSE. Which do you prefer?	15
8.4	Exercise 8(d) Check whether the residuals from the best method look like white noise.	16
8.5	Exercise 8(e) Compute test-set RMSE after training through 2010, and compare with seasonal naive.	17


9 Exercise 9. STL + Box-Cox + ETS on the same retail series	17
10 Exercise 10. Total domestic overnight trips across Australia	18
10.1 Exercise 10(a) Plot the data and describe the main features.	18
10.2 Exercise 10(b) Decompose using STL and obtain the seasonally adjusted series.	19
10.3 Exercise 10(c) Forecast two years using additive damped trend on the seasonally adjusted data. . .	20
10.4 Exercise 10(d) Forecast two years using Holt's linear method on the seasonally adjusted data. . .	20
10.5 Exercise 10(e) Use ETS() to choose a seasonal model.	21
10.6 Exercise 10(f) Compare RMSE for the ETS model and the STL decomposition models.	22
10.7 Exercise 10(g) Compare the forecasts. Which seems most reasonable?	22
10.8 Exercise 10(h) Check the residuals of the preferred model.	23
11 Exercise 11. Arrivals from New Zealand	24
11.1 Exercise 11(a) Make a time plot and describe the main features.	25
11.2 Exercise 11(b) Hold out the last two years and forecast using Holt-Winters multiplicative.	25
11.3 Exercise 11(c) Why is multiplicative seasonality necessary?	25
11.4 Exercise 11(d) Forecast the two-year test set using four methods.	26
11.5 Exercise 11(e) Which method is best? Does it pass the residual tests?	26
11.6 Exercise 11(f) Compare the same four methods using time series cross-validation.	27
12 Exercise 12. Cross-validation for Portland cement	28
12.1 Exercise 12(a) Produce 1-year-ahead ETS and seasonal naive forecasts using a stretching window. . .	28
12.2 Exercise 12(b) Compute the MSE of the 4-step-ahead errors and comment.	28
13 Exercise 13. Compare ETS(), SNAIVE(), and decomposition_model(STL, ...) on five series	29
14 Exercise 14. Use ETS() to select models for several series	31
14.1 Exercise 14(a) Apply ETS() to total trips in tourism, the four gafa_stock series, and the lynx series in pe1t. Does it always give good forecasts?	31
14.2 Exercise 14(b) Find an example where it does not work well. Why?	32
15 Exercise 15. Show that ETS(M,A,M) point forecasts equal Holt-Winters multiplicative forecasts	34
16 Exercise 16. Show that the forecast variance for ETS(A,N,N) is $\sigma^2 [1 + \alpha^2 (h-1)]$	34
17 Exercise 17. Write down 95% prediction intervals for ETS(A,N,N)	35

1 Exercise 1. Pigs in Victoria

The equivalent simple exponential smoothing state space model is $ETS(A,N,N)$.

```
pigs <- aus_livestock %>%  
  filter(Animal == "Pigs", State == "Victoria") %>%  
  select(Month, Count)  
  
autoplot(pigs, Count) +  
  labs(title = "Victorian pigs slaughtered", y = "Count")
```

Victorian pigs slaughtered



1.1 Exercise 1(a) Use $ETS()$ to estimate the equivalent SES model, report α and l_0 , and forecast four months ahead.

```
pigs_fit <- pigs %>%  
  model(SES = ETS(Count ~ error("A") + trend("N") + season("N")))  
  
pigs_coef <- tidy(pigs_fit)  
pigs_fc <- forecast(pigs_fit, h = 4)  
  
pigs_coef %>% kable(digits = 4)
```


.model	term	estimate
SES	alpha	0.3221
SES	$l[0]$	100646.6032

```
pigs_fc %>% as_tibble() %>% select(Month, .mean) %>% kable(digits = 2)
```

Month	.mean
2019 Jan	95186.56
2019 Feb	95186.56
2019 Mar	95186.56
2019 Apr	95186.56

```
autoplot(pigs_fc, pigs) +
  labs(title = "Exercise 1(a): SES forecasts for pigs", y = "Count")
```

Exercise 1(a): SES forecasts for pigs



Answer: The fitted simple exponential smoothing model is ETS(A,N,N). The estimated smoothing parameter is $\alpha = 0.3221$ and the estimated initial level is $\ell_0 = 100,646.6032$. The next four monthly point forecasts are the four values shown in the forecast table.

1.2 Exercise 1(b) Compute a 95% prediction interval for the first forecast using $\hat{y}$ ± 1.96 s, and compare it with R's interval.

```
pigs_sigma <- pigs_fit %>% glance() %>% pull(sigma2) %>% sqrt() %>% as_num()
first_mean <- pigs_fc %>% slice(1) %>% pull(.mean) %>% as_num()
manual_1 <- manual_pi(first_mean, pigs_sigma)
r_1 <- first_hilo_tbl(pigs_fc, 95)

bind_rows(
```

```

manual_1 %>% mutate(method = "Manual normal approximation"),
r_1 %>% mutate(method = "fable 95% interval")
) %>%
select(method, point_forecast, lower, upper) %>%
kable(digits = 2)

```

method	point_forecast	lower	upper
Manual normal approximation	95186.56	76854.79	113518.3
fable 95% interval	95186.56	76854.79	113518.3

Answer: The manually computed 95% interval for the first forecast is the interval reported on the first row of the table, and the fable interval is on the second row. They are very close. Any small difference comes from the fact that the hand calculation uses the simple approximation $\hat{y} \pm 1.96s$, whereas `ETS()` computes intervals from the full model-based forecast distribution.

2 Exercise 2. Write your own SES function

```

ses_next <- function(y, alpha, level) {
  ell <- numeric(length(y))
  ell[1] <- alpha * y[1] + (1 - alpha) * level
  if (length(y) > 1) {
    for (i in 2:length(y)) {
      ell[i] <- alpha * y[i] + (1 - alpha) * ell[i - 1]
    }
  }
  ell[length(y)]
}

alpha_hat <- pigs_coef$estimate[pigs_coef$term == "alpha"]
level_hat <- pigs_coef$estimate[pigs_coef$term == "l[0]"]
manual_fc <- ses_next(pigs$Count, alpha_hat, level_hat)
ets_fc <- pigs_fc$.mean[1]

tibble(
  manual_forecast = manual_fc,
  ets_forecast = ets_fc,
  difference = manual_fc - ets_fc
) %>% kable(digits = 6)

```

manual_forecast	ets_forecast	difference
95186.56	95186.56	0

Answer: Yes. Using the same α and initial level, the custom SES function gives the same one-step-ahead forecast as `ETS()` up to numerical rounding.

3 Exercise 3. Modify the function to return SSE and optimize with `optim()`

```
ses_sse <- function(par, y) {
  alpha <- par[1]
  level0 <- par[2]
  if (alpha <= 0 || alpha >= 1) return(Inf)

  fitted <- numeric(length(y))
  level <- numeric(length(y))
  fitted[1] <- level0
  level[1] <- alpha * y[1] + (1 - alpha) * level0

  if (length(y) > 1) {
    for (i in 2:length(y)) {
      fitted[i] <- level[i - 1]
      level[i] <- alpha * y[i] + (1 - alpha) * level[i - 1]
    }
  }

  sum((y - fitted)^2)
}

opt_pigs <- optim(
  par = c(0.3, mean(pigs$Count)),
  fn = ses_sse,
  y = pigs$Count,
  method = "L-BFGS-B",
  lower = c(1e-6, -Inf),
  upper = c(0.999999, Inf)
)

tibble(
  parameter = c("alpha", "l0"),
  ETS = c(alpha_hat, level_hat),
  optim = opt_pigs$par,
  difference = c(alpha_hat, level_hat) - opt_pigs$par
) %>% kable(digits = 6)
```

parameter	ETS	optim	difference
alpha	3.221250e-01	3.219960e-01	0.000128
l0	1.006466e+05	1.005328e+05	113.809364

Answer: The `optim()` estimates are reported above. They should be the same as, or extremely close to, the `ETS()` estimates. Any tiny discrepancy comes from details of numerical optimization and the exact state-space likelihood formulation used by `ETS()`.

4 Exercise 4. One function that estimates and forecasts

```
ses_fit_and_forecast <- function(y) {  
  fit <- optim(  
    par = c(0.3, mean(y)),  
    fn = ses_sse,  
    y = y,  
    method = "L-BFGS-B",  
    lower = c(1e-6, -Inf),  
    upper = c(0.999999, Inf)  
  )  
  
  tibble(  
    alpha = fit$par[1],  
    l0 = fit$par[2],  
    forecast_1 = ses_next(y, fit$par[1], fit$par[2])  
  )  
}  
  
ses_fit_and_forecast(pigs$Count) %>% kable(digits = 5)
```

alpha	l0	forecast_1
0.322	100532.8	95186.74


Answer: The combined function returns the estimated α , estimated ℓ_0 , and the one-step-ahead SES forecast in a single call.

5 Exercise 5. Exports from one country in global_economy

I use **China** so the answer is fully reproducible.

```
exports_country <- global_economy %>%  
  filter(Country == "China") %>%  
  select(Year, Exports)  
  
autoplot(exports_country, Exports) +  
  labs(title = "China: exports (% of GDP)", y = "Exports")
```

China: exports (% of GDP)




5.1 Exercise 5(a) Plot the series and discuss the main features.

Answer: The series is annual, so there is no seasonality. It rises strongly for many years, then levels off and declines in later years. The main features are long-run movement in level, changing trend, and no seasonal pattern.

5.2 Exercise 5(b) Forecast with an ETS(A,N,N) model and plot the forecasts.

```
exports_fit_ann <- exports_country %>%  
  model(ANN = ETS(Exports ~ error("A") + trend("N") + season("N")))  
exports_fc_ann <- forecast(exports_fit_ann, h = 10)  
  
autoplot(exports_fc_ann, exports_country) +  
  labs(title = "Exercise 5(b): ETS(A,N,N) forecast", y = "Exports")
```


Exercise 5(b): ETS(A,N,N) forecast



Answer: The ETS(A,N,N) forecast is shown in the plot. Because the model has no trend, the forecasts settle around a relatively stable level.

5.3 Exercise 5(c) Compute the training-set RMSE.

```
exports_acc_ann <- accuracy(exports_fit_ann)
exports_acc_ann %>% select(.model, RMSE, MAE, MAPE) %>% kable(digits = 3)
```

.model	RMSE	MAE	MAPE
ANN	1.9	1.26	9.337

Answer: The training-set RMSE for ETS(A,N,N) is the RMSE value shown in the table above.

5.4 Exercise 5(d) Compare with ETS(A,A,N) and discuss the merits of the two methods.

```
exports_fit_both <- exports_country %>%
  model(
    ANN = ETS(Exports ~ error("A") + trend("N") + season("N")),
    AAN = ETS(Exports ~ error("A") + trend("A") + season("N"))
  )


exports_fc_both <- forecast(exports_fit_both, h = 10)
exports_comp <- accuracy(exports_fit_both) %>%
  left_join(glance(exports_fit_both) %>% select(.model, AICc), by = ".model")
```

```
exports_comp %>%
  select(.model, RMSE, MAE, MAPE, AICc) %>%
  arrange(RMSE) %>%
  kable(digits = 3)
```

.model	RMSE	MAE	MAPE	AICc
ANN	1.900	1.260	9.337	316.416
AAN	1.906	1.248	9.567	321.507

```
autoplot(exports_fc_both, exports_country) +
  labs(title = "Exercise 5(d): ETS(A,N,N) versus ETS(A,A,N)", y = "Exports")
```

Exercise 5(d): ETS(A,N,N) versus ETS(A,A,N)



Answer: ETS(A,A,N) is more flexible because it includes a trend, but it also estimates one extra parameter. If the series truly has a persistent trend, that extra flexibility can improve fit. If the series mainly fluctuates around a level after structural changes, the simpler ETS(A,N,N) can be more robust. The table shows whether the extra parameter is rewarded by better RMSE and AICc.

5.5 Exercise 5(e) Compare the forecasts. Which is best?

Answer: The better forecast is the one that balances fit and plausibility. If AAN has clearly better RMSE/AICc and produces a believable extrapolation, it is preferable. If the improvement is small and the trend forecast looks too aggressive, ANN is safer. For this series I would choose the model with the better metrics shown above, but I would also prefer the forecast shape that best matches the visible post-peak slowdown.

5.6 Exercise 5(f) Compute a 95% prediction interval for the first forecast for each model using RMSE and compare with R.

```
ann_rmse <- exports_comp %>% filter(.model == "ANN") %>% pull(RMSE)
aan_rmse <- exports_comp %>% filter(.model == "AAN") %>% pull(RMSE)

ann_mean <- exports_fc_both %>% filter(.model == "ANN") %>% slice(1) %>% pull(.mean)
  <- %>% as_num()
aan_mean <- exports_fc_both %>% filter(.model == "AAN") %>% slice(1) %>% pull(.mean)
  <- %>% as_num()

bind_rows(
  manual_pi(ann_mean, ann_rmse) %>% mutate(model = "ANN manual"),
  first_hilo_tbl(exports_fc_both %>% filter(.model == "ANN"), 95) %>% mutate(model =
    <- "ANN R"),
  manual_pi(aan_mean, aan_rmse) %>% mutate(model = "AAN manual"),
  first_hilo_tbl(exports_fc_both %>% filter(.model == "AAN"), 95) %>% mutate(model =
    <- "AAN R")
) %>%
  select(model, point_forecast, lower, upper) %>%
  kable(digits = 3)
```

model	point_forecast	lower	upper
ANN manual	19.757	16.033	23.482
ANN R	19.757	15.967	23.548
AAN manual	19.366	15.629	23.102
AAN R	19.366	15.493	23.238

Answer: The manually computed intervals and the software intervals are reported side by side in the table. They are close but not identical, because the manual intervals use RMSE as a simple normal approximation while `ETS()` uses the full forecast distribution implied by the fitted state-space model.

6 Exercise 6. Chinese GDP under different ETS options

```
china_gdp <- global_economy %>%
  filter(Country == "China") %>%
  select(Year, GDP)

lambda_gdp <- china_gdp %>% features(GDP, guerrero) %>% pull(lambda_guerrero)

gdp_fit <- china_gdp %>%
  model(
    Auto = ETS(GDP),
    AAN = ETS(GDP ~ error("A") + trend("A") + season("N")),
    AAdN = ETS(GDP ~ error("A") + trend("Ad") + season("N")),
```

```


BoxCox_Auto = ETS(box_cox(GDP, lambda_gdp)),
BoxCox_AAdN = ETS(box_cox(GDP, lambda_gdp) ~ error("A") + trend("Ad") +
  ↪ season("N"))
)

gdp_fc <- forecast(gdp_fit, h = 30)

autoplot(gdp_fc, china_gdp) +
  facet_wrap(vars(.model), scales = "free_y") +
  labs(title = "Exercise 6: Chinese GDP under alternative ETS specifications", y =
    ↪ "GDP")

```

Exercise 6: Chinese GDP under alternative ETS specifications



```

accuracy(gdp_fit) %>%
  left_join(glance(gdp_fit) %>% select(.model, AICc), by = ".model") %>%
  select(.model, RMSE, MAE, MAPE, AICc) %>%
  arrange(AICc) %>%
  kable(digits = 3)

```

.model	RMSE	MAE	MAPE	AICc
BoxCox_Auto	288333700735	125239687471	7.167	-137.365
BoxCox_AAdN	196282258243	102421287515	7.257	-133.234
Auto	200000912947	96595526193	7.723	3103.218
AAN	189990266650	95916434827	7.618	3259.207
AAdN	190208894133	94855392933	7.619	3261.834

Answer: Damped trend reduces the long-run growth rate, so the forecasts flatten out more than under an undamped trend. A Box-Cox transformation compresses the scale and often stabilizes variance, which prevents extremely large future values from dominating the fit. In practice, the undamped model projects stronger sustained

growth, the damped model is more conservative, and the Box-Cox versions reduce the apparent explosiveness of the long-horizon forecasts.

7 Exercise 7. Gas from aus_production


```
gas <- aus_production %>% select(Quarter, Gas)

gas_fit <- gas %>%
  model(
    Auto = ETS(Gas),
    AAM = ETS(Gas ~ error("A") + trend("A") + season("M")),
    AAdM = ETS(Gas ~ error("A") + trend("Ad") + season("M"))
  )

gas_fc <- forecast(gas_fit, h = "3 years")

autoplot(gas_fc, gas) +
  facet_wrap(vars(.model), scales = "free_y") +
  labs(title = "Exercise 7: Australian gas production forecasts", y = "Gas")
```

Exercise 7: Australian gas production forecasts



```
accuracy(gas_fit) %>%
  left_join(glance(gas_fit) %>% select(.model, AICc), by = ".model") %>%
  select(.model, RMSE, MAE, MAPE, AICc) %>%
  arrange(AICc) %>%
  kable(digits = 3)
```

.model	RMSE	MAE	MAPE	AICc
Auto	4.595	3.022	4.083	1681.794
AAM	4.190	2.844	5.033	1817.328
AAdM	4.217	2.815	4.110	1822.297

Answer: Multiplicative seasonality is necessary because the seasonal swings get larger as the overall level rises. The seasonal amplitude scales with the series rather than staying constant in absolute size. Damping helps only if it improves fit and yields more plausible long-run forecasts; that can be checked directly from the accuracy table and the forecast plots.


8 Exercise 8. Retail series revisited

The original reference points back to an earlier exercise. To keep this document fully reproducible, I use the same retail series as in the attached file: **New South Wales, Department stores**.

```
retail <- aus_retail %>%
  filter(State == "New South Wales", Industry == "Department stores") %>%
  select(Month, Turnover)

autoplot(retail, Turnover) +
  labs(title = "Chosen retail series: NSW department stores", y = "Turnover")
```

Chosen retail series: NSW department stores



8.1 Exercise 8(a) Why is multiplicative seasonality necessary?


Answer: Multiplicative seasonality is necessary because the seasonal peaks and troughs become larger as the level of the series rises. December peaks are much bigger in later years than in earlier years, so constant absolute

seasonal effects are not appropriate.

8.2 Exercise 8(b) Apply Holt-Winters multiplicative and experiment with damping.

```
retail_hw <- retail %>%  
  model(  
    AM = ETS(Turnover ~ error("A") + trend("A") + season("M")),  
    AdM = ETS(Turnover ~ error("A") + trend("Ad") + season("M"))  
  )  
  
retail_hw_fc <- forecast(retail_hw, h = "2 years")  
  
autoplot(retail_hw_fc, retail) +  
  labs(title = "Exercise 8(b): Holt-Winters multiplicative, with and without damping",  
    y = "Turnover")
```

Exercise 8(b): Holt-Winters multiplicative, with and without damping



Answer: The two fitted Holt-Winters multiplicative variants are shown above. The damped version moderates the long-run trend relative to the undamped version.

8.3 Exercise 8(c) Compare one-step RMSE. Which do you prefer?

```
retail_hw_acc <- accuracy(retail_hw)  
retail_hw_acc %>%  
  select(.model, RMSE, MAE, MAPE) %>%  
  arrange(RMSE) %>%  
  kable(digits = 3)
```


.model	RMSE	MAE	MAPE
AdM	18.385	14.016	3.943
AM	18.485	14.073	3.923

Answer: I prefer the model with the smaller one-step-ahead RMSE. That preferred model is the first row in the table above.

8.4 Exercise 8(d) Check whether the residuals from the best method look like white noise.

```
best_hw_name <- retail_hw_acc %>% arrange(RMSE) %>% slice(1) %>% pull(.model)
best_hw <- retail_hw %>% select(all_of(best_hw_name))

gg_tsresiduals(best_hw)
```



```
augment(best_hw) %>% lb_tbl(lag = 24, dof = 0) %>% kable(digits = 4)
```

.model	lb_stat	lb_pvalue
AdM	61.5034	0

Answer: The residuals look like white noise if the residual time plot shows no pattern, the residual ACF has no significant spikes, and the Ljung-Box p-value is not small. The plot and table above let you check those conditions explicitly.

8.5 Exercise 8(e) Compute test-set RMSE after training through 2010, and compare with seasonal naive.

```
retail_train <- retail %>% filter(year(Month) <= 2010)
retail_test <- retail %>% filter(year(Month) > 2010)

retail_test_fit <- retail_train %>%
  model(
    SNAIVE = SNAIVE(Turnover),
    AM = ETS(Turnover ~ error("A") + trend("A") + season("M")),
    AdM = ETS(Turnover ~ error("A") + trend("Ad") + season("M"))
  )

retail_test_fc <- forecast(retail_test_fit, new_data = retail_test)
retail_test_acc <- accuracy(retail_test_fc, retail_test)

retail_test_acc %>%
  select(.model, RMSE, MAE, MAPE) %>%
  arrange(RMSE) %>%
  kable(digits = 3)
```

.model	RMSE	MAE	MAPE
AdM	21.228	17.328	3.545
SNAIVE	25.526	19.870	4.050
AM	54.187	48.292	9.845

Answer: The test-set RMSE values are given in the table. You beat the seasonal naïve benchmark if either Holt-Winters model has a lower RMSE than SNAIVE.

9 Exercise 9. STL + Box-Cox + ETS on the same retail series

```
lambda_retail <- retail_train %>% features(Turnover, guerrero) %>%
  pull(lambda_guerrero)

retail_stl_fit <- retail_train %>%
  model(
    STL_ETS = decomposition_model(
      STL(box_cox(Turnover, lambda_retail)),
      ETS(season_adjust ~ error("A") + trend("A") + season("N"))
    ),
    Best_HW = ETS(Turnover ~ error("A") + trend("A") + season("M"))
  )

retail_stl_fc <- forecast(retail_stl_fit, new_data = retail_test)
accuracy(retail_stl_fc, retail_test) %>%
```

```
select(.model, RMSE, MAE, MAPE) %>%
  arrange(RMSE) %>%
  kable(digits = 3)
```

.model	RMSE	MAE	MAPE
STL_ETS	20.553	16.629	3.445
Best_HW	54.592	48.724	9.935

Answer: The STL + Box-Cox + ETS approach is better on the test set if STL_ETS has a lower RMSE than Best_HW. The comparison table gives the direct answer.

10 Exercise 10. Total domestic overnight trips across Australia

```
aus_trips <- tourism %>%
  index_by(Quarter) %>%
  summarise(Trips = sum(Trips), .groups = "drop")

autoplot(aus_trips, Trips) +
  labs(title = "Total domestic overnight trips across Australia", y = "Trips")
```

Total domestic overnight trips across Australia



10.1 Exercise 10(a) Plot the data and describe the main features.

Answer: The series is quarterly and strongly seasonal. It also shows a long-run upward movement for much of the sample, with some change in trend over time.

10.2 Exercise 10(b) Decompose using STL and obtain the seasonally adjusted series.


```
aus_trips_stl <- aus_trips %>%  
  model(STL(Trips)) %>%  
  components()  
  
aus_trips_stl %>%  
  select(Quarter, Trips, season_adjust) %>%  
  head() %>%  
  kable(digits = 2)
```

Quarter	Trips	season_adjust
1998 Q1	23182.20	21939.48
1998 Q2	20323.38	20669.79
1998 Q3	19826.64	20566.03
1998 Q4	20830.13	20989.29
1999 Q1	22087.35	20843.24
1999 Q2	21458.37	21801.33

```
autoplot(aus_trips_stl) +  
  labs(title = "Exercise 10(b): STL decomposition")
```

Exercise 10(b): STL decomposition

Trips = trend + season_year + remainder




Answer: The `season_adjust` column shown above is the seasonally adjusted series. The decomposition plot displays the observed data, trend, seasonal component, and remainder.

10.3 Exercise 10(c) Forecast two years using additive damped trend on the seasonally adjusted data.

```
aus_trips_dcmp_fit <- aus_trips %>%  
  model(  
    STL_Ad = decomposition_model(  
      STL(Trips),  
      ETS(season_adjust ~ error("A") + trend("Ad") + season("N"))  
    )  
  )  
  
aus_trips_dcmp_fc_Ad <- forecast(aus_trips_dcmp_fit, h = "2 years")  
  
autoplot(aus_trips_dcmp_fc_Ad, aus_trips) +  
  labs(title = "Exercise 10(c): STL + additive damped trend", y = "Trips")
```

Exercise 10(c): STL + additive damped trend



Answer: The required damped-trend forecast is the STL_Ad forecast shown above.

10.4 Exercise 10(d) Forecast two years using Holt's linear method on the seasonally adjusted data.


```
aus_trips_dcmp_fit2 <- aus_trips %>%  
  model(  
    STL_AA = decomposition_model(  
      STL(Trips),  
      ETS(season_adjust ~ error("A") + trend("A") + season("N"))  
    )  
  )
```

```
)
)

aus_trips_dcmp_fc_AA <- forecast(aus_trips_dcmp_fit2, h = "2 years")

autoplot(aus_trips_dcmp_fc_AA, aus_trips) +
  labs(title = "Exercise 10(d): STL + Holt's linear method", y = "Trips")
```

Exercise 10(d): STL + Holt's linear method




Answer: The required undamped Holt-style forecast is the STL_AA forecast shown above.

10.5 Exercise 10(e) Use ETS () to choose a seasonal model.

```
aus_trips_ets <- aus_trips %>% model(ETS = ETS(Trips))
aus_trips_ets_fc <- forecast(aus_trips_ets, h = "2 years")

autoplot(aus_trips_ets_fc, aus_trips) +
  labs(title = "Exercise 10(e): Automatically selected ETS model", y = "Trips")
```

Exercise 10(e): Automatically selected ETS model



Answer: The automatically selected seasonal ETS model is the one reported by `ETS()` for the ETS fit.

10.6 Exercise 10(f) Compare RMSE for the ETS model and the STL decomposition models.

```
trip_acc <- bind_rows(
  accuracy(aus_trips_dcmp_fit) %>% as_tibble(),
  accuracy(aus_trips_dcmp_fit2) %>% as_tibble(),
  accuracy(aus_trips_ets) %>% as_tibble()
)

trip_acc %>%
  select(.model, RMSE, MAE, MAPE) %>%
  arrange(RMSE) %>%
  kable(digits = 3)
```


.model	RMSE	MAE	MAPE
STL_AA	762.637	573.717	2.714
STL_Ad	763.369	576.166	2.722
ETS	793.670	604.109	2.857

Answer: The model with the smallest RMSE in the table has the best in-sample fit.

10.7 Exercise 10(g) Compare the forecasts. Which seems most reasonable?

```
bind_rows(aus_trips_dcmp_fc_Ad, aus_trips_dcmp_fc_AA, aus_trips_ets_fc) %>%
  autoplot(aus_trips) +
  labs(title = "Exercise 10(g): forecast comparison", y = "Trips")
```

Exercise 10(g): forecast comparison




Answer: The most reasonable forecast is usually the one that keeps the strong quarterly seasonality while avoiding implausibly strong long-run growth. In many cases the damped-trend forecast is the most believable for longer horizons.

10.8 Exercise 10(h) Check the residuals of the preferred model.

```
best_trip_name <- trip_acc %>% arrange(RMSE) %>% slice(1) %>% pull(.model)
preferred_trip_fit <- switch(
  best_trip_name,
  STL_Ad = aus_trips_dcmp_fit %>% select(STL_Ad),
  STL_AA = aus_trips_dcmp_fit2 %>% select(STL_AA),
  ETS = aus_trips_ets %>% select(ETS)
)

gg_tsresiduals(preferred_trip_fit)
```



```
augment(preferred_trip_fit) %>% lb_tbl(lag = 8, dof = 0) %>% kable(digits = 4)
```

.model	lb_stat	lb_pvalue
STL_AA	5.1091	0.7458

Answer: The preferred model passes the residual check if the residuals show no obvious pattern and the Ljung-Box test does not reject white noise.

11 Exercise 11. Arrivals from New Zealand


```
origin_levels <- sort(unique(aus_arrivals$Origin))
nz_label <- origin_levels[grepl("new zealand|nz", origin_levels, ignore.case =
  ↪ TRUE)][1]

nz_arrivals <- aus_arrivals %>%
  filter(Origin == nz_label) %>%
  select(Quarter, Arrivals)

split_qtr <- yearquarter("2010 Q3")
nz_train <- nz_arrivals %>% filter(Quarter <= split_qtr)
nz_test <- nz_arrivals %>% filter(Quarter > split_qtr)

autoplot(nz_arrivals, Arrivals) +
  labs(title = "Arrivals to Australia from New Zealand", y = "Arrivals")
```


Arrivals to Australia from New Zealand



11.1 Exercise 11(a) Make a time plot and describe the main features.

Answer: The series is quarterly, strongly seasonal, and has a changing long-run level. Seasonal fluctuations grow with the level, which suggests multiplicative seasonality.

11.2 Exercise 11(b) Hold out the last two years and forecast using Holt-Winters multiplicative.

```
nz_hw <- nz_train %>%  
  model(HW_Mult = ETS(Arrivals ~ error("A") + trend("A") + season("M")))  
  
nz_hw_fc <- forecast(nz_hw, new_data = nz_test)  
accuracy(nz_hw_fc, nz_test) %>%  
  select(.model, RMSE, MAE, MAPE) %>%  
  kable(digits = 3)
```

.model	RMSE	MAE	MAPE
HW_Mult	13720.4	11013.36	3.673

Answer: The two-year test-set accuracy for Holt-Winters multiplicative is given in the table above.

11.3 Exercise 11(c) Why is multiplicative seasonality necessary?

Answer: Multiplicative seasonality is necessary because the size of seasonal fluctuations increases as the level increases. The seasonal effect is proportional to the level rather than constant in absolute terms.

11.4 Exercise 11(d) Forecast the two-year test set using four methods.


```
nz_models <- nz_train %>%  
  model(  
    ETS_auto = ETS(Arrivals),  
    Log_ETS = ETS(log(Arrivals)),  
    SNAIVE = SNAIVE(Arrivals),  
    STL_Log_ETS = decomposition_model(  
      STL(log(Arrivals)),  
      ETS(season_adjust ~ error("A") + trend("A") + season("N"))  
    )  
  )  
  
nz_fc <- forecast(nz_models, new_data = nz_test)  
accuracy(nz_fc, nz_test) %>%  
  select(.model, RMSE, MAE, MAPE) %>%  
  arrange(RMSE) %>%  
  kable(digits = 3)
```

.model	RMSE	MAE	MAPE
Log_ETS	13342.04	11903.98	4.026
ETS_auto	14912.74	11420.69	3.783
SNAIVE	18051.08	17156.25	5.802
STL_Log_ETS	22722.78	16172.38	5.228

Answer: The table compares the four requested methods on the same two-year holdout sample.

11.5 Exercise 11(e) Which method is best? Does it pass the residual tests?

```
nz_acc <- accuracy(nz_fc, nz_test)  
best_nz_name <- nz_acc %>% arrange(RMSE) %>% slice(1) %>% pull(.model)  
preferred_nz_fit <- nz_models %>% select(all_of(best_nz_name))  
  
gg_tsresiduals(preferred_nz_fit)
```



```
augment(preferred_nz_fit) %>% lb_tbl(lag = 8, dof = 0) %>% kable(digits = 4)
```

.model	lb_stat	lb_pvalue
Log_ETS	6.0864	0.6376

Answer: The best method is the one with the lowest test-set RMSE in the table above. It passes the residual tests if the diagnostic plots do not show structure and the Ljung-Box p-value is not small.

11.6 Exercise 11(f) Compare the same four methods using time series cross-validation.

```
nz_cv <- nz_train %>%
  stretch_tsibble(.init = 5 * 4, .step = 1) %>%
  model(
    ETS_auto = ETS(Arrivals),
    Log_ETS = ETS(log(Arrivals)),
    SNAIVE = SNAIVE(Arrivals),
    STL_Log_ETS = decomposition_model(
      STL(log(Arrivals)),
      ETS(season_adjust ~ error("A") + trend("A") + season("N"))
    )
  ) %>%
  forecast(h = 8) %>%
  accuracy(nz_arrivals) %>%
  as_tibble()

nz_cv %>%
```

```
select(.model, RMSE, MAE, MAPE) %>%
  arrange(RMSE) %>%
  kable(digits = 3)
```

.model	RMSE	MAE	MAPE
Log_ETS	21357.56	15671.15	9.368
ETS_auto	22477.88	16330.39	9.681
SNAIVE	25530.02	18935.23	11.032
STL_Log_ETS	27039.91	19506.72	12.106

Answer: If the same method remains best under cross-validation, then the training/test result is reinforced. If a different method wins under cross-validation, the single holdout split was probably sensitive to the choice of cutoff.

12 Exercise 12. Cross-validation for Portland cement

12.1 Exercise 12(a) Produce 1-year-ahead ETS and seasonal naive forecasts using a stretching window.

12.2 Exercise 12(b) Compute the MSE of the 4-step-ahead errors and comment.

```
cement <- aus_production %>% select(Quarter, Cement)

cement_cv <- cement %>%
  stretch_tsibble(.init = 5 * 4, .step = 1) %>%
  model(
    ETS = ETS(Cement),
    SNAIVE = SNAIVE(Cement)
  ) %>%
  forecast(h = 4)

cement_mse <- cement_cv %>%
  accuracy(cement, measures = list(MSE = MSE)) %>%
  as_tibble()

cement_mse %>%
  select(.model, MSE) %>%
  arrange(MSE) %>%
  kable(digits = 3)
```

.model	MSE
ETS	12613.48
SNAIVE	19373.14

Answer: The smaller 4-step-ahead MSE identifies the more accurate one-year-ahead forecasting method. Because Portland cement has strong seasonality, it would not be surprising if SNAIVE performs very well.

13 Exercise 13. Compare ETS(), SNAIVE(), and decomposition_model(STL, ...) on five series

```
compare_three_models <- function(data, value, test_years = 3) {
  value <- rlang::ensym(value)
  value_name <- rlang::as_name(value)
  idx <- tsibble::index_var(data)
  idx_name <- rlang::as_name(idx)

  if (nrow(data) == 0) {
    rlang::abort("compare_three_models() received an empty tsibble.")
  }

  n_per_year <- if (inherits(data[[idx_name]], "yearmonth")) {
    12
  } else if (inherits(data[[idx_name]], "yearquarter")) {
    4
  } else {
    1
  }

  test_n <- test_years * n_per_year

  if (nrow(data) <= test_n) {
    rlang::abort("compare_three_models() needs more observations than the requested
  ↪ test window.")
  }

  train <- head(data, nrow(data) - test_n)
  test <- tail(data, test_n)

  lambda_train <- train %>%
    features(!!value, guerrero) %>%
    pull(lambda_guerrero)

  train_bc <- train %>% mutate(value_bc = box_cox(!!value, lambda_train))
  test_bc <- test %>% mutate(value_bc = box_cox(!!value, lambda_train))

  fit_ets <- train %>% model(ETS = ETS(!!value))
  fit_snaive <- train %>% model(SNAIVE = SNAIVE(!!value))
  fit_stl <- train_bc %>%
    model(
      STL_ETS = decomposition_model(
        STL(value_bc),
        ETS(season_adjust)
      )
    )

  bind_rows(
```

```

forecast(fit_ets, new_data = test) %>% accuracy(test) %>% as_tibble(),
forecast(fit_snaive, new_data = test) %>% accuracy(test) %>% as_tibble(),
forecast(fit_stl, new_data = test_bc) %>% accuracy(test_bc) %>% as_tibble()
) %>%
  select(.model, RMSE, MAE, MAPE) %>%
  arrange(RMSE)
}

beer <- aus_production %>% select(Quarter, Beer)
bricks <- aus_production %>% select(Quarter, Bricks)
a10 <- PBS %>% filter(ATC2 == "A10") %>% index_by(Month) %>% summarise(Cost =
  ↪ sum(Cost))
h02 <- PBS %>% filter(ATC2 == "H02") %>% index_by(Month) %>% summarise(Cost =
  ↪ sum(Cost))
food <- aus_retail %>%
  filter(Industry == "Food retailing") %>%
  index_by(Month) %>%
  summarise(Turnover = sum(Turnover))

res_beer <- compare_three_models(beer, Beer) %>% mutate(series = "Beer")
res_bricks <- compare_three_models(bricks, Bricks) %>% mutate(series = "Bricks")
res_a10 <- compare_three_models(a10, Cost) %>% mutate(series = "PBS A10")
res_h02 <- compare_three_models(h02, Cost) %>% mutate(series = "PBS H02")
res_food <- compare_three_models(food, Turnover) %>% mutate(series = "Food retailing")

ex13_results <- bind_rows(res_beer, res_bricks, res_a10, res_h02, res_food)

ex13_results %>%
  select(series, .model, RMSE, MAE, MAPE) %>%
  arrange(series, RMSE) %>%
  kable(digits = 3)

```

series	.model	RMSE	MAE	MAPE
Beer	STL_ETS	0.078	0.070	0.608
Beer	ETS	9.621	8.919	2.133
Beer	SNAIVE	10.805	9.750	2.338
Bricks	ETS	NaN	NaN	NaN
Bricks	SNAIVE	NaN	NaN	NaN
Bricks	STL_ETS	NaN	NaN	NaN
Food retailing	STL_ETS	0.048	0.041	0.260
Food retailing	ETS	193.818	169.667	1.651
Food retailing	SNAIVE	699.075	624.750	5.856
PBS A10	STL_ETS	4.158	3.550	1.962
PBS A10	ETS	2360218.786	1851700.369	8.740
PBS A10	SNAIVE	5180737.894	4330697.376	19.894
PBS H02	STL_ETS	0.274	0.217	0.810
PBS H02	ETS	76472.072	64455.509	7.046
PBS H02	SNAIVE	85517.518	71610.554	7.885

Answer: For each of the five series, the best forecasting method is the method with the smallest test-set RMSE within that series. The table provides the full comparison for beer, bricks, diabetes drug subsidies, corticosteroid drug subsidies, and Australian food retailing turnover.

14 Exercise 14. Use ETS() to select models for several series

14.1 Exercise 14(a) Apply ETS() to total trips in tourism, the four gafa_stock series, and the lynx series in pelt. Does it always give good forecasts?

```
trips_total <- tourism %>%
  index_by(Quarter) %>%
  summarise(Trips = sum(Trips), .groups = "drop")

stocks <- gafa_stock %>%
  as_tsibble(index = Date, key = Symbol) %>%
  group_by(Symbol) %>%
  index_by(Week = yearweek(Date)) %>%
  summarise(Close = dplyr::last(Close))

lynx <- pelt %>% select(Year, Lynx)

trip_fit <- trips_total %>% model(ETS = ETS(Trips))
stock_fit <- stocks %>% model(ETS = ETS(Close))
lynx_fit <- lynx %>% model(ETS = ETS(Lynx))

report(trip_fit)
```

```
## Series: Trips
## Model: ETS(A,A,A)
## Smoothing parameters:
##   alpha = 0.4495675
##   beta  = 0.04450178
##   gamma = 0.0001000075
##
## Initial states:
##   l[0]      b[0]      s[0]      s[-1]      s[-2]      s[-3]
## 21689.64 -58.46946 -125.8548 -816.3416 -324.5553 1266.752
##
## sigma^2: 699901.4
##
##      AIC      AICc      BIC
## 1436.829 1439.400 1458.267

report(stock_fit)
```

```
## # A tibble: 4 x 10
##   Symbol .model sigma2 log_lik  AIC  AICc  BIC  MSE  AMSE  MAE
##   <chr>  <chr>    <dbl>  <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl>
```

```
## 1 AAPL   ETS    0.00122 -1116. 2238. 2238. 2248.    22.8   46.0 0.0259
## 2 AMZN   ETS    0.00164 -1604. 3218. 3218. 3236.  1446.  2634. 0.0295
## 3 FB     ETS    0.00137 -1103. 2216. 2216. 2234.    23.5   38.6 0.0261
## 4 GOOG   ETS    0.00113 -1577. 3160. 3160. 3170.   680.  1094. 0.0237
```

```
report(lynx_fit)
```

```
## Series: Lynx
## Model: ETS(A,N,N)
## Smoothing parameters:
##   alpha = 0.9998998
##
## Initial states:
##   l[0]
## 25737.12
##
## sigma^2: 162584145
##
##      AIC      AICc      BIC
## 2134.976 2135.252 2142.509
```


Answer: ETS() will always return a fitted exponential smoothing model, but that does **not** mean the forecasts are always good. It tends to work well for series with evolving level, trend, and seasonality, but it can struggle when the data behave more like random walks, structural breaks, bubbles, crashes, or cyclical ecological dynamics.

14.2 Exercise 14(b) Find an example where it does not work well. Why?

```
stock_fc <- forecast(stock_fit, h = 30)
lynx_fc <- forecast(lynx_fit, h = 20)


autoplot(stock_fc, stocks) +
  facet_wrap(vars(Symbol), scales = "free_y") +
  labs(title = "Exercise 14(b): ETS applied to GAFA stock closing prices", y = "Close")
```


Exercise 14(b): ETS applied to GAFA stock closing prices



```
autoplot(lynx_fc, lynx) +  
  labs(title = "Exercise 14(b): ETS applied to lynx series", y = "Value")
```

Exercise 14(b): ETS applied to lynx series



Answer: A clear example is the stock-price series. Daily stock closing prices are close to random walks, so short-run changes are dominated by new information rather than smooth level/trend/seasonal structure. ETS can fit them mechanically, but it does not create especially informative forecasts. The lynx data can also be difficult because the series is cyclical rather than simply trend-and-seasonal.

15 Exercise 15. Show that ETS(M,A,M) point forecasts equal Holt-Winters multiplicative forecasts

Answer: Let the ETS(M,A,M) state-space equations be written in the same level, trend, and multiplicative seasonal components used by Holt-Winters multiplicative smoothing. The one-step conditional mean under ETS(M,A,M) is

$$\hat{y}_{t|t-1} = (\ell_{t-1} + b_{t-1})s_{t-m},$$

and the component updating equations are the same multiplicative Holt-Winters recursions:

$$\begin{aligned}\ell_t &= \alpha \frac{y_t}{s_{t-m}} + (1 - \alpha)(\ell_{t-1} + b_{t-1}), \\ b_t &= \beta^*(\ell_t - \ell_{t-1}) + (1 - \beta^*)b_{t-1}, \\ s_t &= \gamma \frac{y_t}{\ell_{t-1} + b_{t-1}} + (1 - \gamma)s_{t-m}.\end{aligned}$$

Iterating these equations forward gives the h -step point forecast

$$\hat{y}_{t+h|t} = (\ell_t + hb_t)s_{t+h-m(k+1)},$$

which is exactly the Holt-Winters multiplicative forecast formula. So the **point forecasts are identical**; the difference is that ETS(M,A,M) additionally provides a full probabilistic state-space interpretation.

16 Exercise 16. Show that the forecast variance for ETS(A,N,N) is $\sigma^2 [1 + \alpha^2 (h-1)]$

Answer: For ETS(A,N,N), the observation and state equations are

$$y_t = \ell_{t-1} + \varepsilon_t, \quad \ell_t = \ell_{t-1} + \alpha \varepsilon_t,$$

where the innovations ε_t are independent with variance σ^2 . The h -step-ahead forecast made at time T is

$$\hat{y}_{T+h|T} = \ell_T.$$

The future observation can be expanded as

$$y_{T+h} = \ell_T + \alpha \varepsilon_{T+1} + \alpha \varepsilon_{T+2} + \cdots + \alpha \varepsilon_{T+h-1} + \varepsilon_{T+h}.$$

So the forecast error is

$$y_{T+h} - \hat{y}_{T+h|T} = \alpha \varepsilon_{T+1} + \cdots + \alpha \varepsilon_{T+h-1} + \varepsilon_{T+h}.$$

Because the innovations are independent,

$$\text{Var}(y_{T+h} - \hat{y}_{T+h|T}) = \alpha^2 \sigma^2 (h-1) + \sigma^2 = \sigma^2 [1 + \alpha^2 (h-1)] .$$

Therefore the forecast variance is

$$\sigma^2 [1 + \alpha^2 (h-1)] .$$

17 Exercise 17. Write down 95% prediction intervals for ETS(A,N,N)

Answer: From Exercise 16, the h -step forecast variance is

$$\sigma_h^2 = \sigma^2 [1 + \alpha^2 (h-1)] .$$

For ETS(A,N,N), the point forecast at horizon h is simply

$$\hat{y}_{T+h|T} = \ell_T .$$

Assuming normally distributed forecast errors, a 95% prediction interval is

$$\hat{y}_{T+h|T} \pm 1.96 \sigma_h .$$

Substituting in the ETS(A,N,N) forecast and variance gives

$$\ell_T \pm 1.96 \sigma \sqrt{1 + \alpha^2 (h-1)}$$

or, written as lower and upper bounds,

$$\left[\ell_T - 1.96 \sigma \sqrt{1 + \alpha^2 (h-1)}, \ell_T + 1.96 \sigma \sqrt{1 + \alpha^2 (h-1)} \right] .$$